

PostgreSQL Page Internals

- PostgreSQL Page Representation
- TOAST'ing

❖ PostgreSQL Page Representation

Functions: **src/backend/storage/page/*.c**

Definitions: **src/include/storage/bufpage.h**

Each page is 8KB (default **BLCKSZ**) and contains:

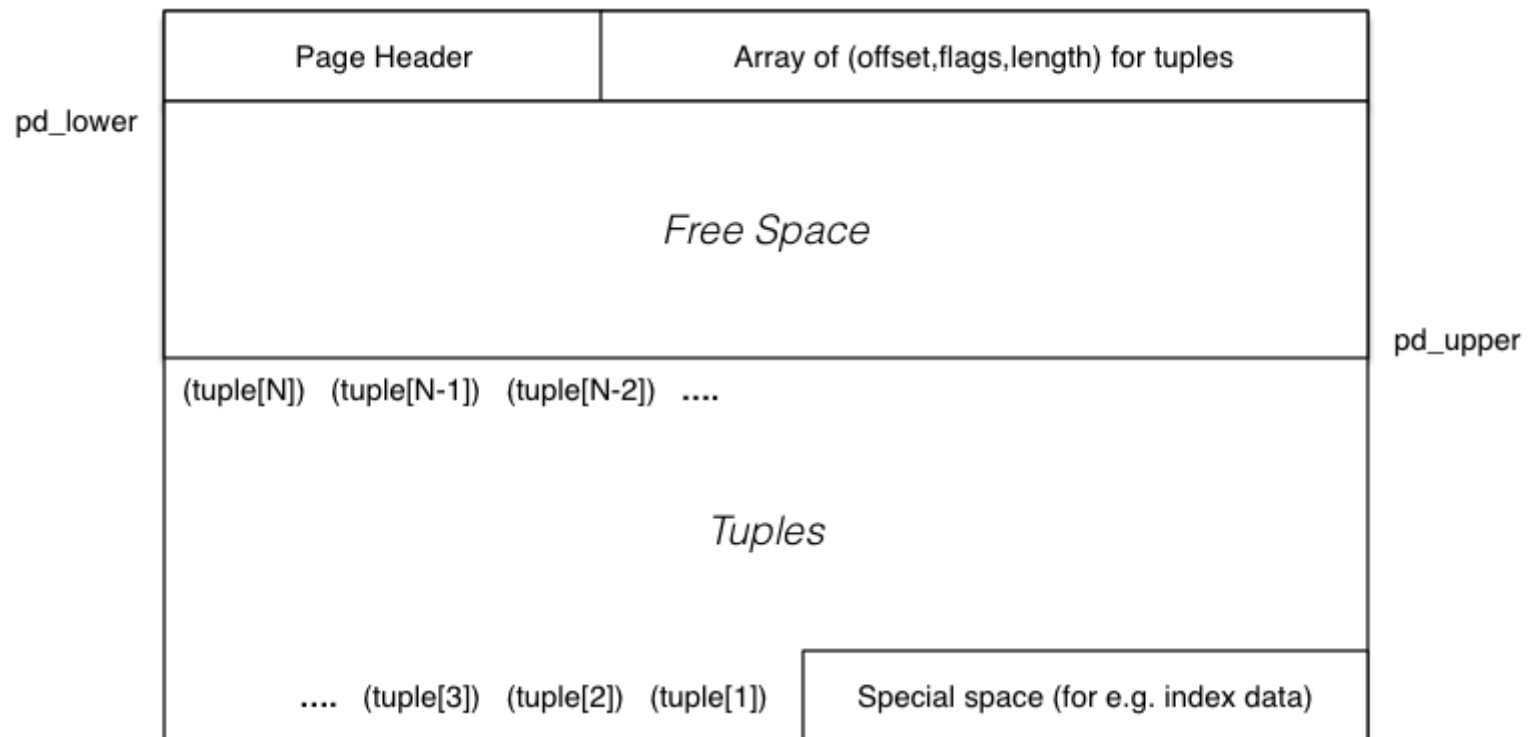
- header (free space pointers, flags, xact data)
- array of (offset,length) pairs for tuples in page
- free space region (between array and tuple data)
- actual tuples themselves (inserted from end towards start)
- (optionally) region for special data (e.g. index data)

Large data items are stored in separate (TOAST) files (implicit)

Also supports ~SQL-standard BLOBs (explicit large data items)

❖ PostgreSQL Page Representation (cont)

PostgreSQL page layout:



❖ PostgreSQL Page Representation (cont)

Page-related data types:

```
// a Page is simply a pointer to start of buffer
typedef Pointer Page;

// indexes into the tuple directory
typedef uint16 LocationIndex;

// entries in tuple directory (line pointer array)
typedef struct ItemIdData
{
    unsigned    lp_off:15,    // tuple offset from start of page
                lp_flags:2,   // unused,normal,redirect,dead
                lp_len:15;    // length of tuple (bytes)
} ItemIdData;
```

❖ PostgreSQL Page Representation (cont)

Page-related data types: (cont)

```
typedef struct PageHeaderData
{
    XLogRecPtr      pd_lsn;          // xact log record for last change
    uint16          pd_tli;          // xact log reference information
    uint16          pd_flags;        // flag bits (e.g. free, full, ...
    LocationIndex   pd_lower;        // offset to start of free space
    LocationIndex   pd_upper;        // offset to end of free space
    LocationIndex   pd_special;      // offset to start of special space
    uint16          pd_pagesize_version;
    TransactionId   pd_prune_xid;    // is pruning useful in data page?
    ItemIdData      pd_linp[1];     // beginning of line pointer array
} PageHeaderData;

typedef PageHeaderData *PageHeader;
```

❖ PostgreSQL Page Representation (cont)

Operations on **Pages**:

void PageInit(Page page, Size pageSize, ...)

- initialize a **Page** buffer to empty page
- in particular, sets **pd_lower** and **pd_upper**

**OffsetNumber PageAddItem(Page page,
Item item, Size size, ...)**

- insert one tuple (or index entry) into a **Page**
- fails if: not enough free space, too many tuples

void PageRepairFragmentation(Page page)

- compact tuple storage to give one large free space region

❖ PostgreSQL Page Representation (cont)

PostgreSQL has two kinds of pages:

- **heap pages** which contain tuples
- **index pages** which contain index entries

Both kinds of page have the same page layout.


One important difference:

- index entries tend to be smaller than tuples
- can typically fit more index entries per page

❖ TOAST'ing

TOAST = The Oversized-Attribute Storage Technique

- handles storage of large attribute values (> 2KB) (e.g. long **text**)



❖ TOAST'ing (cont)

Large attribute values are stored out-of-line (i.e. in separate file)

- "value" of attribute in tuple is a reference to TOAST data
- TOAST'd values may be compressed
- TOAST'd values are stored in 2K chunks

Strategies for storing TOAST-able columns ...

- **PLAIN** ... allows no compression or out-of-line storage
- **EXTENDED** ... allows both compression and out-of-line storage
- **EXTERNAL** ... allows out-of-line storage but not compression
- **MAIN** ... allows compression but not out-of-line storage

Produced: 23 Feb 2021