

# COMP9315 Week 2 Sessions

---

- Things to Note
- Exercise: Relation Scan Cost
- Exercise: PostgreSQL Files
- Exercise: Buffer Cost Benefit (i)
- Exercise: Buffer Cost Benefit (ii)
- Exercise: Clock-sweep Page Replacement
- Exercise: Fixed-length Records
- Exercise: Inserting/Deleting Fixed-length Records
- Exercise: Inserting Variable-length Records
- Exercise: PostgreSQL Pages
- Exercise: Space Utilisation

## ❖ Things to Note

---

### Course Notes

- more detailed than Slides ... need upgrading ... coming soon

### Quiz 1

- due before Friday 23:59 ... so far 90/470 submissions

### Dropping/Enrolling

- if you drop COMP9315, will need help enrolling in replacement

### Unix skills

- Home Computing playlist on  
<https://www.youtube.com/channel/UCi3Kf5eONlwV6QgNHiYqVzg>

## ❖ Exercise: Relation Scan Cost

Consider a table  $R(x,y,z)$  with  $10^5$  tuples, implemented as

- number of tuples  $r = 10,000$
- average size of tuples  $R = 200$  bytes
- size of data pages  $B = 4096$  bytes
- time to read one data page  $T_r = 10\text{msec}$
- time to check one tuple  $1\text{ usec}$
- time to form one result tuple  $1\text{ usec}$
- time to write one result page  $T_r = 10\text{msec}$

Calculate the total time-cost for answering the query:

```
insert into S select * from R where x > 10;
```

if 50% of the tuples satisfy the condition.

COMP9315 21T1 ◇ Week 2 Sessions ◇ [2/11]

## ❖ Exercise: PostgreSQL Files

---

In your PostgreSQL server

- examine the content of the **\$PGDATA/base** directory
- find the directory containing the **pizza** database
- find the file in this directory for the **People** table
- examine the contents of the **People** file
- what are the other files in the directory?
- are there **forks** in any of your databases?

## ❖ Exercise: Buffer Cost Benefit (i)

Assume that:

- the **Customer** relation has  $b_C$  pages (e.g. 10)
- the **Employee** relation has  $b_E$  pages (e.g. 4)

Compute how many page reads occur ...

- if we have only 2 buffers (i.e. effectively no buffer pool)
- if we have 20 buffers
- when a buffer pool with MRU replacement strategy is used
- when a buffer pool with LRU replacement strategy is used

For the last two, buffer pool has  $n=3$  slots ( $n < b_C$  and  $n < b_E$ )

## ❖ Exercise: Buffer Cost Benefit (ii)

---

If the tables were larger, the above analysis would be tedious.

Write a C program to simulate buffer pool usage

- assuming a nested loop join as above
- **argv[1]** gives number of pages in "outer" table
- **argv[2]** gives number of pages in "inner" table
- **argv[3]** gives number of slots in buffer pool
- **argv[4]** gives replacement strategy (LRU,MRU,FIFO-Q)

## ❖ Exercise: Clock-sweep Page Replacement

Using the following data type for buffer frame descriptors:

```
struct FrameDesc {  
    Tag pid;        // ID of page in frame e.g. "R0"  
    int pin;        // number tx's using this page  
    int usage;      // clock-sweep usage counter  
}
```

Show how the buffer pool changes for

- $n = 4$ ,  $b_R = 3$ ,  $b_S = 4$ ,  $b_T = 6$
- when executing **select \* from T** via sequential scan
- when executing **select \* from R join S** using nested-loop join



## ❖ Exercise: Fixed-length Records

---

Give examples of table definitions

- which result in fixed-length records
- which result in variable-length records

```
create table R ( ... );
```

What are the common features of each type of table?

## ❖ Exercise: Inserting/Deleting Fixed-length Records

---

For each of the following Page formats:

- compacted/packed free space
- unpacked free space (with bitmap)

Implement

- a suitable data structure to represent a **Page**
- a function to insert a new record
- a function to delete a record

## ❖ Exercise: Inserting Variable-length Records

For both of the following page formats

1. variable-length records, with compacted free space
2. variable-length records, with fragmented free space

implement the **insert()** function.

Use the above page format, but also assume:

- page size is 1024 bytes
- tuples start on 4-byte boundaries
- references into page are all 8-bits (1 byte) long
- a function **recSize(r)** gives size in bytes

## ❖ Exercise: PostgreSQL Pages

---

Draw diagrams of a PostgreSQL heap page

- when it is initially empty
- after three tuples have been inserted with lengths of 60, 80, and 70 bytes
- after the 80 byte tuple is deleted (but before vacuuming)
- after a new 50 byte tuple is added

Show the values in the tuple header.

Assume that there is no special space in the page.

## ❖ Exercise: Space Utilisation

Consider the following page/record information:

- page size = 1KB = 1024 bytes =  $2^{10}$  bytes
- records: **(a:int, b:varchar(20), c:char(10), d:int)**
- records are all aligned on 4-byte boundaries
- **c** field padded to ensure **d** starts on 4-byte boundary
- each record has 4 field-offsets at start of record (each 1 byte)
- **char(10)** field rounded up to 12-bytes to preserve alignment
- maximum size of **b** values = 20 bytes; average size = 16 bytes
- page has 32-bytes of header information, starting at byte 0
- only insertions, no deletions or updates

Calculate  $c$  = average number of records per page.

Produced: 25 Feb 2021