


Tuple Representation

- Tuples
- Records vs Tuples
- Converting Records to Tuples
- Operations on Records
- Operations on Tuples
- Fixed-length Records
- Variable-length Records
- Data Types
- Field Descriptors

❖ Tuples

Each **page** contains a collection of **tuples**



What do tuples contain? How are they structured internally?

❖ Records vs Tuples

A **table** is defined by a **schema**, e.g.

```
create table Employee (  
    id    integer primary key,  
    name  varchar(20) not null,  
    job   varchar(10),  
    dept  smallint references Dept(id)  
);
```

where a schema is a collection of attributes (name,type,constraints)

Reminder: schema information (meta-data) is also stored, in the DB catalog

❖ Records vs Tuples (cont)

Tuple = collection of attribute values based on a schema, e.g.

(33357462, 'Neil Young', 'Musician', 277)

iid:integer name:varchar(20) job:varchar(10) dept: smallint

Record = sequence of bytes, containing data for one tuple, e.g.

01101001	11001100	01010101	00111100	10100011	01011111	01011010	
----------	----------	----------	----------	----------	----------	----------	--

Bytes need to be interpreted relative to schema to get tuple

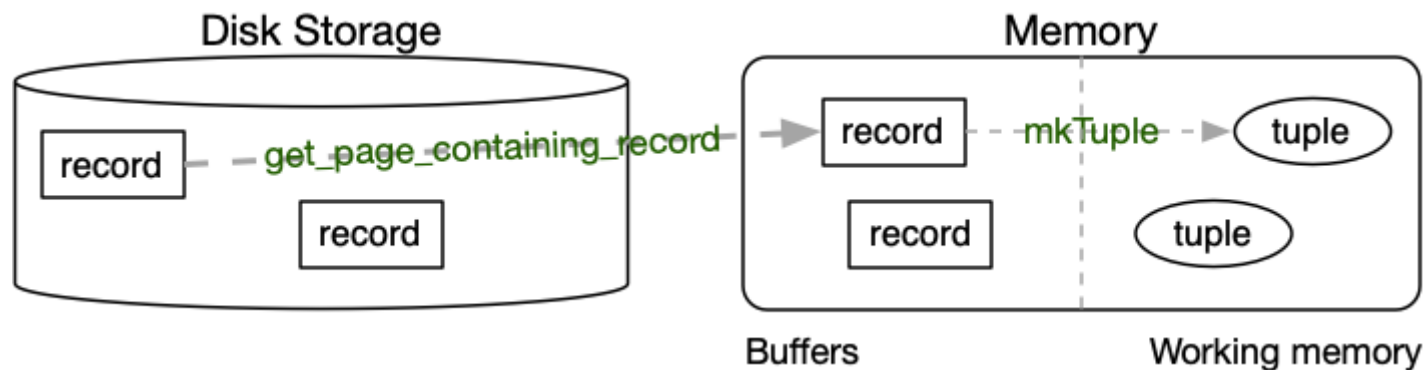
❖ Converting Records to Tuples

A **Record** is an array of bytes (**byte[]**)

- representing the data values from a typed **Tuple**
- stored on disk (persistent) or in a memory buffer

A **Tuple** is a collection of named, typed values (cf. C **struct**)

- to manipulate the values, need an "interpretable" structure
- stored in working memory, and temporary



❖ Converting Records to Tuples (cont)

Information on how to interpret bytes in a record ...

- may be contained in schema data in DBMS catalog
- may be stored in the page directory
- may be stored in the record (in a record header)
- may be stored partly in the record and partly in the schema

For variable-length records, some formatting info ...

- must be stored in the record or in the page directory
- at the least, need to know how many bytes in each varlen value

❖ Operations on Records

Common operation on records ... access record via **RecordId**:

```
Record get_record(Relation rel, RecordId rid) {  
    (pid,tid) = rid;  
    Page buf = get_page(rel, pid);  
    return get_bytes(rel, buf, tid);  
}
```

Cannot use a **Record** directly; need a **Tuple**:

```
Relation rel = ... // relation schema  
Record rec = get_record(rel, rid)  
Tuple t = mkTuple(rel, rec)
```

Once we have a **Tuple**, we can access individual attributes/fields

❖ Operations on Tuples

Once we have a record, we need to interpret it as a tuple ...

```
Tuple t = mkTuple(rel, rec)
```

- convert record to tuple data structure for relation **rel**

Once we have a tuple, we want to examine its contents ...

```
Typ getTypField(Tuple t, int i)
```

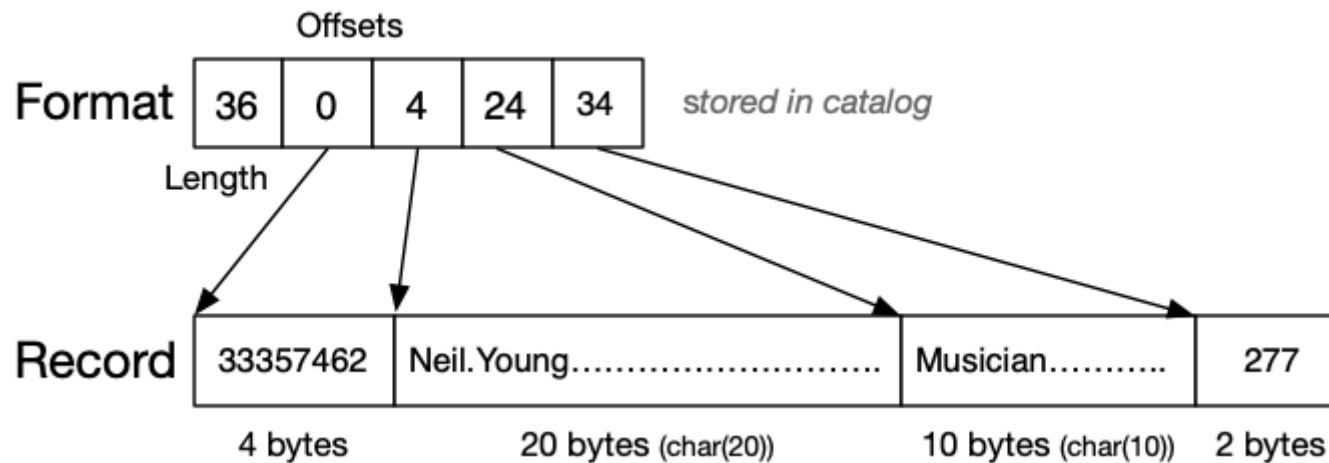
- extract the **i**'th field from a **Tuple** as a value of type *Typ*

E.g. **int** x = getIntField(t,1), **char** *s = getStrField(t,2)

❖ Fixed-length Records

A possible encoding scheme for fixed-length records:

- record format (length + offsets) stored in catalog
- data values stored in fixed-size slots in data pages



Since record format is frequently used at query time, cache in memory.

❖ Variable-length Records

Possible encoding schemes for variable-length records:

- Prefix each field by length

4	33357462	10	Neil Young	8	Musician	2	277
---	----------	----	------------	---	----------	---	-----

- Terminate fields by delimiter

33357462	X	Neil Young	X	Musician	X	277	X
----------	---	------------	---	----------	---	-----	---

- Array of offsets

len | offsets[] | data

34	10	14	24	32	33357462	Neil Young	Musician	277
----	----	----	----	----	----------	------------	----------	-----

❖ Data Types

DBMSs typically define a fixed set of base types, e.g.

DATE, FLOAT, INTEGER, NUMBER(*n*), VARCHAR(*n*), ...

This determines implementation-level data types for field values:

DATE	time_t
FLOAT	float, double
INTEGER	int, long
NUMBER(<i>n</i>)	int[] (?)
VARCHAR(<i>n</i>)	char[]

PostgreSQL allows new base types to be added

❖ Field Descriptors

A **Tuple** could be implemented as

- a list of field descriptors for a record instance
(where a **FieldDesc** gives (offset,length,type) information)
- along with a reference to the **Record** data

```
typedef struct {  
    ushort    nfields;    // number of fields/attrs  
    ushort    data_off;   // offset in struct for data  
    FieldDesc fields[];   // field descriptions  
    Record    data;       // pointer to record in buffer  
} Tuple;
```

Fields are derived from relation descriptor + record instance data.

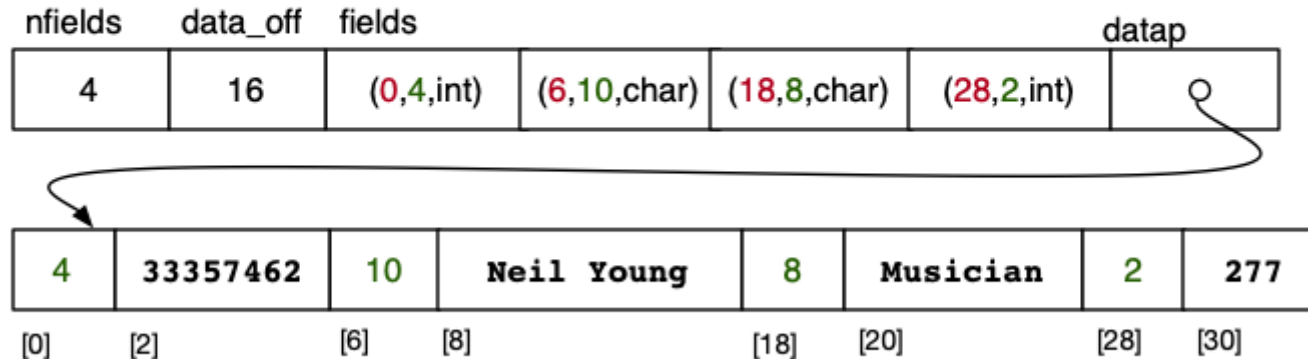
❖ Field Descriptors (cont)

Tuple **data** could be

- a pointer to bytes stored elsewhere in memory



e.g.



Note that the *offset* refers to the length field at the start of each attribute.

❖ Field Descriptors (cont)

Or, tuple **data** could be ...

- appended to **Tuple struct** (used widely in PostgreSQL)

nfields,data_off,fields,data

....	tuple data
------	------------

e.g.

nfields	data_off	fields				
4	16	(0,4,int)	(6,10,char)	(18,8,char)	(28,2,int)	...

...	4	33357462	10	Neil Young	8	Musician	2	277
-----	---	----------	----	------------	---	----------	---	-----

Produced: 27 Feb 2021