

Scanning


- Scanning
- Selection via Scanning
- Iterators
- Example Query
- **next_tuple()** Function
- Relation Copying
- Scanning in PostgreSQL
- Scanning in other File Structures

❖ Scanning

Consider executing the query:

```
select * from Rel;
```

where the relation has a file structure like:



This would be done by a simple scan of all records/tuples.

❖ Scanning (cont)

Abstract view of how the scan might be implemented:

```
for each tuple T in relation Rel {  
    add tuple T to result set  
}
```


Operational view:

```
for each page P in file of relation Rel {  
    for each tuple T in page P {  
        add tuple T to result set  
    }  
}
```

Cost = read every data page once = b

❖ Scanning (cont)

Consider a file with overflow pages, e.g.



❖ Scanning (cont)

In this case, the implementation changes to:

```
for each page P in data file of relation Rel {  
    for each tuple t in page P {  
        add tuple t to result set  
    }  
    for each overflow page V of page P {  
        for each tuple t in page V {  
            add tuple t to result set  
        }  
    }  
}
```

Cost: read each data page and each overflow page once

$$Cost = b + b_{OV}$$


where b_{OV} = total number of overflow pages

❖ Selection via Scanning

Consider a *one* query like:

```
select * from Employee where id = 762288;
```

In an unordered file, search for matching tuple requires:



Guaranteed at most one answer; but could be in any page.

❖ Selection via Scanning (cont)

Overview of scan process:

```
for each page P in relation Employee {  
    for each tuple t in page P {  
        if (t.id == 762288) return t  
    }  
}
```

Cost analysis for *one* searching in unordered file

- best case: read one page, find tuple
- worst case: read all b pages, find in last (or don't find)
- average case: read half of the pages ($b/2$)

Page Costs: $Cost_{avg} = b/2$ $Cost_{min} = 1$ $Cost_{max} = b$

❖ Iterators

Access methods typically involve **iterators**, e.g.

Scan s = start_scan(Relation r, ...)

- commence a scan of relation **r**
- **Scan** may include condition to implement **WHERE**-clause
- **Scan** holds data on progress through file (e.g. current page)

Tuple next_tuple(Scan s)

- return **Tuple** immediately following last accessed one
- returns **NULL** if no more **Tuples** left in the relation

❖ Example Query

Example: simple scan of a table ...

```
select name from Employee
```

implemented as:

```
DB db = openDatabase("myDB");
Relation r = openRelation(db, "Employee", READ);
Scan s = start_scan(r);
Tuple t; // current tuple
while ((t = next_tuple(s)) != NULL) {
    char *name = getStrField(t, 2);
    printf("%s\n", name);
}
```

❖ `next_tuple()` Function

Consider the following possible **Scan** data structure

```
typedef ScanData *Scan;

typedef struct {
    Relation rel;
    Page      *curPage;    // Page buffer
    int        curPID;      // current pid
    int        curTID;      // current tid
} ScanData;
```

Assume tuples are indexed $0..n\text{**Tuples**}(p)-1$

Assume pages are indexed $0..n\text{**Pages**}(rel)-1$

❖ `next_tuple()` Function (cont)

Implementation of **Tuple** `next_tuple(Scan)` function

```
Tuple next_tuple(Scan s)
{
    if (s->curTID >= nTuples(s->page)-1) {
        // get a new page; exhausted current page
        s->curPID++;
        if (s->curPID >= nPages(s->rel))
            return NULL;
        else {
            s->page = get_page(s->rel, s->curPID);
            s->curTID = -1;
        }
    }
    s->curTID++;
    return get_tuple(s->rel, s->page, s->curTID);
}
```

❖ Relation Copying

Consider an SQL statement like:

```
create table T as (select * from S);
```

Effectively, copies data from one table to a new table.

Process:

```
make empty relation T
s = start scan of S
while (t = next_tuple(s)) {
    insert tuple t into relation T
}
```

❖ Relation Copying (cont)

It is possible that **T** is smaller than **S**

- may be unused free space in **S** where tuples were removed
- if **T** is built by simple append, will be compact



❖ Relation Copying (cont)

In terms of existing relation/page/tuple operations:

```
Relation in;           // relation handle (incl. files)
Relation out;          // relation handle (incl. files)
int ipid,opid,tid;     // page and record indexes
Record rec;           // current record (tuple)
Page ibuf,obuf;       // input/output file buffers

in = openRelation("S", READ);
out = openRelation("T", NEW|WRITE);
clear(obuf);  opid = 0;
for (ipid = 0; ipid < nPages(in); ipid++) {
    ibuf = get_page(in, ipid);
    for (tid = 0; tid < nTuples(ibuf); tid++) {
        rec = get_record(ibuf, tid);
        if (!hasSpace(obuf,rec)) {
            put_page(out, opid++, obuf);
            clear(obuf);
        }
        insert_record(obuf,rec);
    }
}
if (nTuples(obuf) > 0) put_page(out, opid, obuf);
```

❖ Scanning in PostgreSQL

Scanning defined in: `backend/access/heap/heapam.c`

Implements iterator data/operations:

- **HeapScanDesc** ... struct containing iteration state
- **scan = heap_beginscan(rel, ..., nkeys, keys)**
- **tup = heap_getnext(scan, direction)**
- **heap_endscan(scan)** ... frees up **scan** struct
- **res = HeapKeyTest(tuple, ..., nkeys, keys)**
... performs **ScanKeys** tests on tuple ... is it a result tuple?

❖ Scanning in PostgreSQL (cont)

```
typedef HeapScanDescData *HeapScanDesc;  
  
typedef struct HeapScanDescData  
{  
    // scan parameters  
    Relation      rs_rd;           // heap relation descriptor  
    Snapshot      rs_snapshot;     // snapshot ... tuple visibility  
    int           rs_nkeys;        // number of scan keys  
    ScanKey       rs_key;          // array of scan key descriptors  
    ...  
    // state set up at initscan time  
    PageNumber    rs_npages;       // number of pages to scan  
    PageNumber    rs_startpage;    // page # to start at  
    ...  
    // scan current state, initially set to invalid  
    HeapTupleData rs_ctup;         // current tuple in scan  
    PageNumber    rs_cpage;        // current page # in scan  
    Buffer         rs_cbuf;         // current buffer in scan  
    ...  
} HeapScanDescData;
```


❖ Scanning in other File Structures

Above examples are for **heap** files

- simple, unordered, maybe indexed, no hashing

Other access file structures in PostgreSQL:

- **btree, hash, gist, gin**
- each implements:
 - startscan, getnext, endscan
 - insert, delete (update=delete+insert)
 - other file-specific operators

Produced: 28 Feb 2021