

Sorting

- The Sort Operation
- Two-way Merge Sort
- Comparison for Sorting
- Cost of Two-way Merge Sort
- n-Way Merge Sort
- Cost of n-Way Merge Sort
- Sorting in PostgreSQL

❖ The Sort Operation

Sorting is explicit in queries only in the **order by** clause

```
select * from Students order by name;
```

Sorting is used internally in other operations:

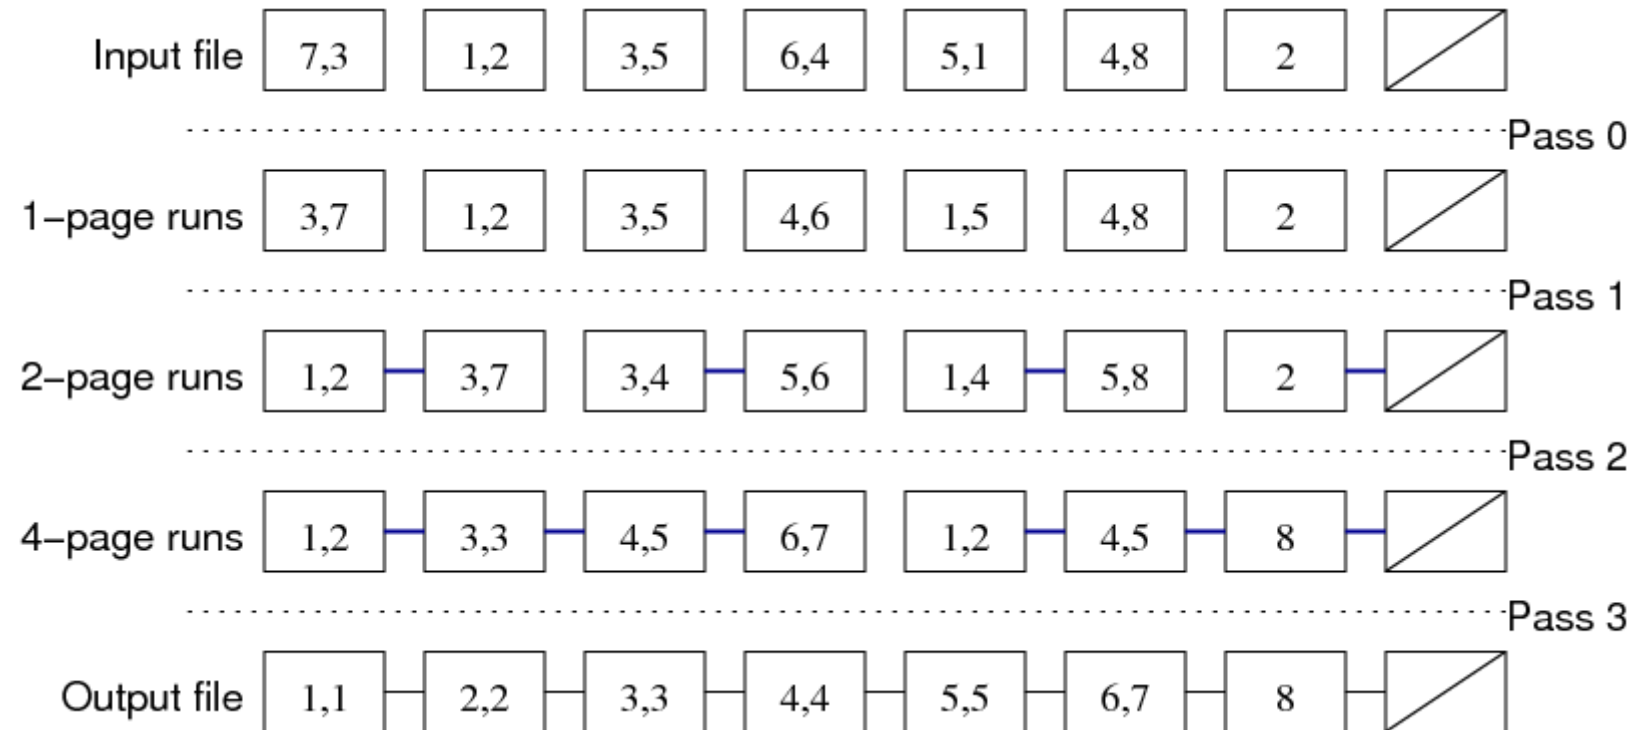
- eliminating duplicate tuples for projection
- ordering files to enhance select efficiency
- implementing various styles of join
- forming tuple groups in **group by**

Sort methods such as quicksort are designed for in-memory data.

For large data on disks, need external sorts such as [merge sort](#).


❖ Two-way Merge Sort

Example:



❖ Two-way Merge Sort (cont)

Requires at least three in-memory buffers:



Assumption: cost of **Merge** operation on two in-memory buffers ≈ 0 .

❖ Comparison for Sorting

Above assumes that we have a function to compare tuples.

Needs to understand ordering on different data types.

Need a function **tupCompare(r1, r2, f)** (cf. C's **strcmp**)

```
int tupCompare(r1, r2, f)
{
    if (r1.f < r2.f) return -1;
    if (r1.f > r2.f) return 1;
    return 0;
}
```

Assume =, <, > are available for all attribute types.

❖ Comparison for Sorting (cont)

In reality, need to sort on multiple attributes and ASC/DESC, e.g.

```
-- example multi-attribute sort
select * from Students
order by age desc, year_enrolled
```

Sketch of multi-attribute sorting function

```
int tupCompare(r1,r2,criteria)
{
    foreach (f,ord) in criteria {
        if (ord == ASC) {
            if (r1.f < r2.f) return -1;
            if (r1.f > r2.f) return 1;
        }
        else {
            if (r1.f > r2.f) return -1;
            if (r1.f < r2.f) return 1;
        }
    }
}
```

```
    }  
    return 0;  
}
```

COMP9315 21T1 ◇ Sorting ◇ [5/14]

❖ Cost of Two-way Merge Sort

For a file containing b data pages:

- require $\text{ceil}(\log_2 b)$ passes to sort,
- each pass requires b page reads, b page writes

Gives total cost: $2.b.\text{ceil}(\log_2 b)$


Example: Relation with $r=10^5$ and $c=50 \Rightarrow b=2000$ pages.

Number of passes for sort: $\text{ceil}(\log_2 2000) = 11$

Reads/writes entire file 11 times! Can we do better?

❖ n-Way Merge Sort


Initial pass uses: B total buffers



Reads B pages at a time, sorts in memory, writes out in order

❖ n-Way Merge Sort (cont)

Merge passes use: $n = B-1$ input buffers, 1 output buffer



❖ n-Way Merge Sort (cont)

Method:

```
// Produce B-page-long runs
for each group of B pages in Rel {
    read B pages into memory buffers
    sort group in memory
    write B pages out to Temp
}
// Merge runs until everything sorted
numberOfRuns = ceil(b/B)
while (numberOfRuns > 1) {
    // n-way merge, where n=B-1
    for each group of n runs in Temp {
        merge into a single run via input buffers
        write run to newTemp via output buffer
    }
    numberOfRuns = ceil(numberOfRuns/n)
    Temp = newTemp // swap input/output files
}
```

❖ Cost of n-Way Merge Sort

Consider file where $b = 4096$, $B = 16$ total buffers:

- pass 0 produces 256×16 -page sorted runs
- pass 1
 - performs 15-way merge of groups of 16-page sorted runs
 - produces 18×240 -page sorted runs (17 full runs, 1 short run)
- pass 2
 - performs 15-way merge of groups of 240-page sorted runs
 - produces 2×3600 -page sorted runs (1 full run, 1 short run)
- pass 3
 - performs 15-way merge of groups of 3600-page sorted runs
 - produces 1×4096 -page sorted runs

(cf. two-way merge sort which needs 11 passes)

❖ Cost of n-Way Merge Sort (cont)

Generalising from previous example ...

For b data pages and B buffers

- first pass: read/writes b pages, gives $b_0 = \text{ceil}(b/B)$ runs
- then need $\text{ceil}(\log_n b_0)$ passes until sorted, where $n = B-1$
- each pass reads and writes b pages (i.e. $2.b$ page accesses)

$Cost = 2.b.(1 + \text{ceil}(\log_n b_0))$, where b_0 and n are defined above

❖ Sorting in PostgreSQL

Sort uses a merge-sort (from Knuth) similar to above:

- `backend/utils/sort/tuplesort.c`
- `include/utils/sortsupport.h`

Tuples are mapped to **SortTuple** structs for sorting:

- containing pointer to tuple and sort key
- no need to reference actual Tuples during sort
- unless multiple attributes used in sort

If all data fits into memory, sort using **qsort()**.

If memory fills while reading, form "runs" and do disk-based sort.

❖ Sorting in PostgreSQL (cont)

Disk-based sort has phases:

- divide input into sorted runs using HeapSort
- merge using N buffers, one output buffer
- N = as many buffers as **workMem** allows

Described in terms of "tapes" ("tape" \cong sorted run)

Implementation of "tapes": [backend/utils/sort/logtape.c](https://github.com/postgres/postgres/blob/master/backend/utils/sort/logtape.c)

❖ Sorting in PostgreSQL (cont)

Sorting comparison operators are obtained via catalog (in ***Type.o***):

```
// gets pointer to function via pg_operator
struct Tuplesortstate { ... SortTupleComparator ... };

// returns negative, zero, positive
ApplySortComparator(Datum datum1, bool isnull1,
                    Datum datum2, bool isnull2,
                    SortSupport sort_helper);
```

Flags in **SortSupport** indicate: ascending/descending, nulls-first/last.

ApplySortComparator() is PostgreSQL's version of **tupCompare()**

Produced: 1 Mar 2021