

COMP9315 Week 3 Exercises

- FYI
- Exercise: Building Tuples in PostgreSQL
- Exercise: How big is a **FieldDesc**?
- Exercise: Cost of Search in Hashed File
- Exercise: Update Operation Costs
- Exercise: Cost of n-Way Merge Sort
- Exercise: Cost of Relation Copy



Quiz 1 marks are now available

- normally available on Monday after due date

Assignment 1

- make a new base data type (like **integer**)
- define: parser/printer, storage structure, operations
- new deadline **9pm** Friday 19 March

Unix skills

- Home Computing playlist on

<https://www.youtube.com/channel/UCi3Kf5eONlwV6QgNHiYqVzg>

❖ Exercise: Building Tuples in PostgreSQL

Examine the code for

```
HeapTuple heap_form_tuple(desc, values[], isnull[])
```

and determine how a PostgreSQL tuple is built

```
HeapTuple = -> HeapTupleData
HeapTupleData =
    (t_len, t_self, t_tableOid, t_data->HeapTupleHeaderData)
HeapTupleHeaderData =
    (t_heap, t_ctid, t_infomask2, t_infomask, t_off, t_bits[], ...)
HeapTupleFields =
    (t_xmin, t_xmax, (t_cid | t_xvac)
TupleDesc =
    (natts, tdtypeid, tdtypmod, tdhasoid, constraints[], atts[])
FormData_pg_attribute =
    (attrelid, attname, atttypid, attlen, attndims, attnotnull, ...)
```

❖ Exercise: How big is a `FieldDesc`?

FieldDesc = (offset,length,type), where

- offset = offset of field within record data
- length = length (in bytes) of field
- type = data type of field

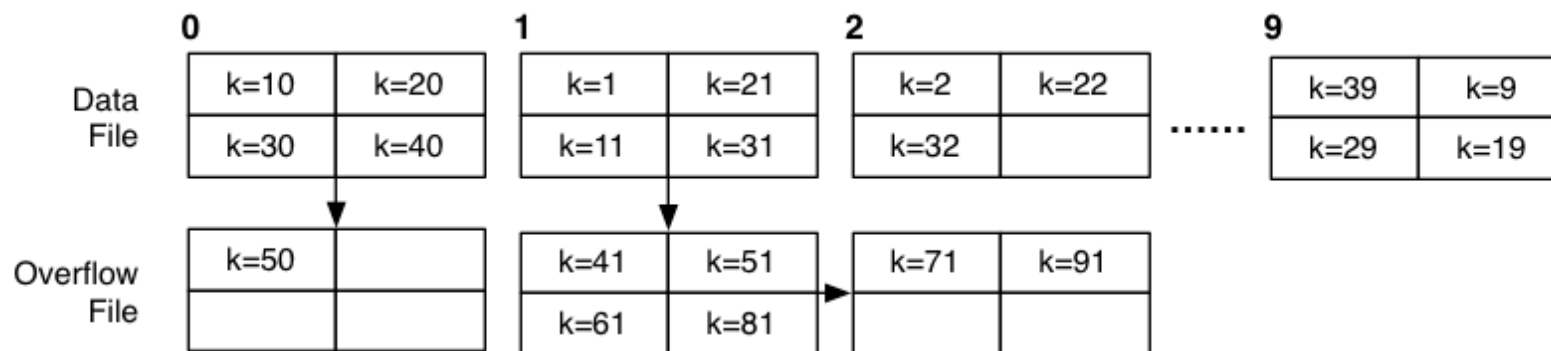
If pages are 8KB in size, how many bits are needed for each?

E.g.

nfields	data_off	fields = FieldDesc[4]			
4	16	(0,4,int)	(6,10,char)	(18,8,char)	(28,2,int)

❖ Exercise: Cost of Search in Hashed File

Consider the hashed file structure $b = 10, c = 4, h(k) = k \% 10$



Describe how the following queries

```
select * from R where k = 51;
```

```
select * from R where k > 50;
```

might be solved in a file structure like the above ($h(k) = k \% b$).

Estimate the minimum and maximum cost (as #pages read)

❖ Exercise: Update Operation Costs

For each of the following file structures

- heap file, sorted file, hash file

Determine #page-reads + #page-writes for insert and delete

You can assume the existence of a file header containing

- values for r , R , b , B , c
- index of first page with free space (and a free list)

Assume also

- each page contains a header and directory as well as tuples
- no buffering (worst case scenario)

❖ Exercise: Cost of n-Way Merge Sort

How many reads+writes to sort the following:

- $r = 1048576$ tuples (2^{20})
- $R = 62$ bytes per tuple (fixed-size)
- $B = 4096$ bytes per page
- $H = 96$ bytes of header data per page
- $D = 1$ presence bit per tuple in page directory
- all pages are full

Consider for the cases:

- 9 total buffers, 8 input buffers, 1 output buffer
- 33 total buffers, 32 input buffers, 1 output buffer
- 257 total buffers, 256 input buffers, 1 output buffer

❖ Exercise: Cost of Relation Copy

Analyse cost for relation copying:

1. if both input and output are heap files
2. if input is sorted and output is heap file
3. if input is heap file and output is sorted

Assume ...

- r records in input file, c records/page
- b_{in} = number of pages in input file
- some pages in input file are *not* full
- all pages in output file are full (except the last)

Give cost in terms of #pages read + #pages written

Produced: 2 Mar 2021