

# Implementing Projection

---

- The Projection Operation
- Sort-based Projection
- Cost of Sort-based Projection
- Hash-based Projection
- Cost of Hash-based Projection
- Projection on Primary Key
- Index-only Projection
- Comparison of Projection Methods
- Projection in PostgreSQL

## ❖ The Projection Operation

Consider the query:

```
select distinct name,age from Employee;
```

If the **Employee** relation has four tuples such as:

```
(94002, John, Sales, Manager, 32)
(95212, Jane, Admin, Manager, 39)
(96341, John, Admin, Secretary, 32)
(91234, Jane, Admin, Secretary, 21)
```

then the result of the projection is:

```
(Jane, 21)    (Jane, 39)    (John, 32)
```

Note that duplicate tuples (e.g. **(John, 32)**) are eliminated.

## ❖ The Projection Operation (cont)

Relies on function **`Tuple projTuple(AttrList, Tuple)`**

- first arg is list of attributes
- second arg is a tuple containing those attributes (and more)
- return value is a new tuple containing only those attributes

Examples, using tuples of type **`(id:int, name:text, degree:int)`**

```
projTuple([id], (1234, 'John', 3778))  
  returns (id=1234)
```

```
projTuple([name, degree], (1234, 'John', 3778))  
  returns (name='John', degree=3778)
```

## ❖ The Projection Operation (cont)

Without **distinct**, projection is straightforward

```
// attrs = [attr1, attr2, ...]
bR = nPages(Rel)
for i in 0 .. bR-1 {
    P = read page i
    for j in 0 .. nTuples(P)-1 {
        T = getTuple(P, j)
        T' = projTuple(attrs, T)
        if (outBuf is full) write and clear
        append T' to outBuf
    }
}
if (nTuples(outBuf) > 0) write
```

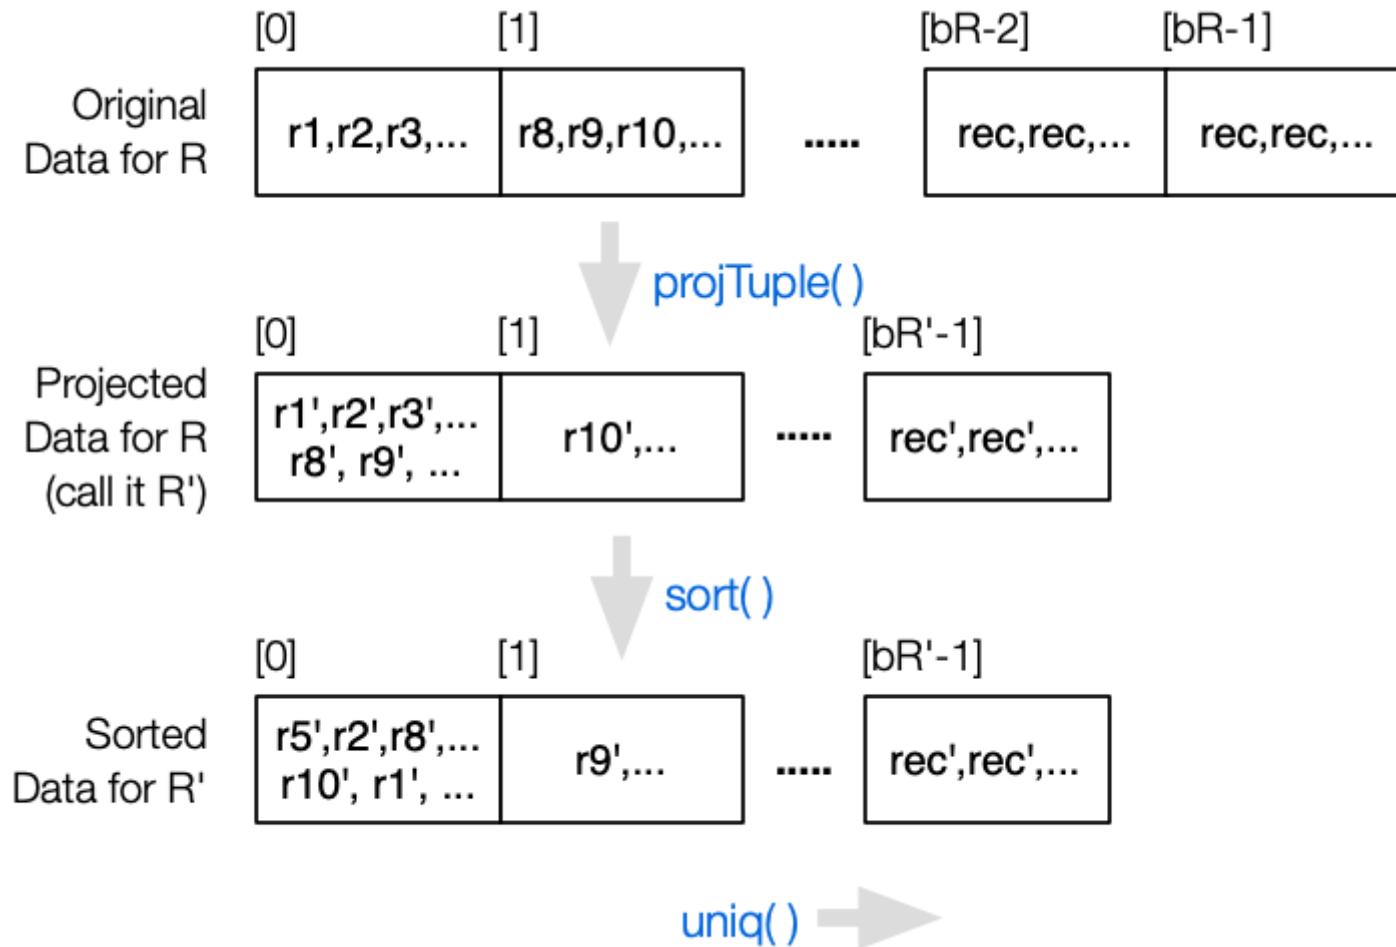
Typically,  $b_{\text{OutFile}} < b_{\text{InFile}}$  (same number of tuples, but tuples are smaller)

## ❖ The Projection Operation (cont)

With **distinct**, the projection operation needs to:

1. scan the entire relation as input
  - already seen how to do scanning
2. create output tuples containing only requested attributes
  - implementation depends on tuple internal structure
  - essentially, make a new tuple with fewer attributes and where the values may be computed from existing attributes
3. eliminate any duplicates produced
  - two approaches: sorting or hashing

## ❖ Sort-based Projection



## ❖ Sort-based Projection (cont)

Requires a temporary file/relation.

```
for each tuple T in RelFile {  
    T' = projTuple([attr1,attr2,...],T)  
    add T' to TempFile  
}
```

```
sort TempFile on [attrs]
```

```
for each tuple T in TempFile {  
    if (T == Prev) continue  
    write T to Result  
    Prev = T  
}
```

Reminder: "**for each tuple**" means page-by-page, tuple-by-tuple

## ❖ Cost of Sort-based Projection

The costs involved are (assuming  $B=n+1$  buffers for sort):


- scanning original relation **Rel**:  $b_R$  (with  $c_R$ )
- writing **Temp** relation:  $b_T$  (smaller tuples,  $c_T > c_R$ , sorted)
- sorting **Temp** relation:  
 $2.b_T.\text{ceil}(\log_n b_0)$  where  $b_0 = \text{ceil}(b_T/B)$
- scanning **Temp**, removing duplicates:  $b_T$
- writing the result relation:  $b_{Out}$  (maybe less tuples)

Cost = sum of above =  $b_R + b_T + 2.b_T.\text{ceil}(\log_n b_0) + b_T + b_{Out}$




## ❖ Hash-based Projection

Partitioning phase:



## ❖ Hash-based Projection (cont)

Duplicate elimination phase:



## ❖ Hash-based Projection (cont)

Algorithm for both phases:

```
for each tuple T in relation Rel {
    T' = mkTuple(attrs,T)
    H = h1(T', n)
    B = buffer for partition[H]
    if (B full) write and clear B
    insert T' into B
}
for each partition P in 0..n-1 {
    for each tuple T in partition P {
        H = h2(T, n)
        B = buffer for hash value H
        if (T not in B) insert T into B
        // assumes B never gets full
    }
    write and clear all buffers
}
```

Reminder: "**for each tuple**" means page-by-page, tuple-by-tuple

COMP9315 21T1 ◇ Implementing Projection ◇ [10/15]

## ❖ Cost of Hash-based Projection

The total cost is the sum of the following:

- scanning original relation **R**:  $b_R$
- writing partitions:  $b_P$  ( $b_R$  vs  $b_P$ ?)
- re-reading partitions:  $b_P$
- writing the result relation:  $b_{Out}$

$$\text{Cost} = b_R + 2b_P + b_{Out}$$

To ensure that  $n$  is larger than the largest partition ...

- use hash functions (h1,h2) with uniform spread
- allocate at least  $\sqrt{b_R} + 1$  buffers
- if insufficient buffers, significant re-reading overhead

## ❖ Projection on Primary Key

No duplicates, so simple approach from above works:

```
bR = nPages(Rel)
for i in 0 .. bR-1 {
    P = read page i
    for j in 0 .. nTuples(P) {
        T = getTuple(P,j)
        T' = projTuple([pk], T)
        if (outBuf is full) write and clear
        append T' to outBuf
    }
}
if (nTuples(outBuf) > 0) write
```

## ❖ Index-only Projection

Can do projection without accessing data file iff ...

- relation is indexed on  $(A_1, A_2, \dots, A_n)$  (indexes described later)
- projected attributes are a prefix of  $(A_1, A_2, \dots, A_n)$

Basic idea:

- scan through index file (which is already sorted on attributes)
- duplicates are already adjacent in index, so easy to skip

Cost analysis ...

- index has  $b_i$  pages (where  $b_i \ll b_R$ )
- Cost =  $b_i$  reads +  $b_{Out}$  writes

## ❖ Comparison of Projection Methods

Difficult to compare, since they make different assumptions:

- index-only: needs an appropriate index
- hash-based: needs buffers and good hash functions
- sort-based: needs only buffers  $\Rightarrow$  use as default

Best case scenario for each (assuming  $n+1$  in-memory buffers):

- index-only:  $b_i + b_{Out} \ll b_R + b_{Out}$
- hash-based:  $b_R + 2.b_P + b_{Out}$
- sort-based:  $b_R + b_T + 2.b_T.\text{ceil}(\log_n b_0) + b_T + b_{Out}$

We normally omit  $b_{Out}$ ... each method produces the same result



## ❖ Projection in PostgreSQL

Code for projection forms part of execution iterators:

- `include/nodes/execnodes.h`
- `backend/executor/execQual.c`

Types:

- **`ProjectionInfo { type, pi_state, pi_exprContext }`**
- **`ExprState { tag, flags, resnull, resvalue, ... }`**

Functions:

- **`ExecProject(projInfo, ...)`** ... extracts projected data
- **`check_sql_fn_retval(...)`** ... evaluates attribute value

Produced: 6 Mar 2021