>>

Selection Overview

- Varieties of Selection
- Implementing Select Efficiently

/\ >>

❖ Varieties of Selection

Selection: **`select * from R where C`**

- filters a subset of tuples from one relation **`R`**
- based on a condition **`C`** on the attribute values

We consider three distinct styles of selection:

- 1-d (one dimensional) (condition uses only *1* attribute)
- *n*-d (multi-dimensional) (condition uses *>1* attribute)
- similarity (approximate matching, with ranking)

Each style has several possible file-structures/techniques.

<< ^ >>

❖ Varieties of Selection (cont)

Selection returns a subset of tuples from a table

- r_q = number of tuples that match query *q*
- b_q = number of pages containing tuples that match query *q*

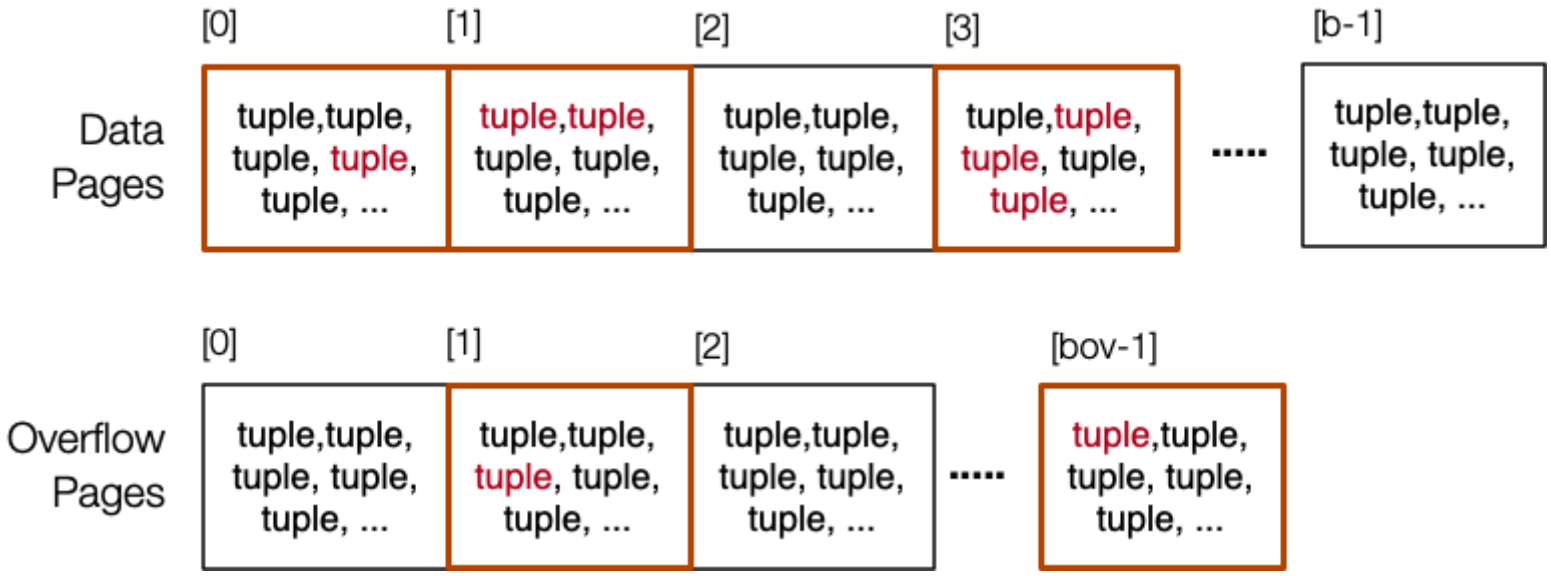

In the diagram, $r_q = 8$, $b_q = 5$

<< ^ >>

❖ Varieties of Selection (cont)

Different categories of selection queries:

one ... queries with at most 1 result ... $0 \leqslant r_q \leqslant 1$, $0 \leqslant b_q \leqslant 1$

- typically, equality test on primary key attribute, e.g.
- **`select * from R where id = 1234`**

pmr ... partial match retrieval ... $0 \leqslant r_q \leqslant r$, $0 \leqslant b_q \leqslant b+b_{ov}$

- conjunction of equality tests on multiple attributes, e.g.
- **`select * from R where age=65`** (1-d)
- **`select * from R where age=65 and gender='m'`** (n-d)

<< ^ >>

❖ Varieties of Selection (cont)

More categories of selection queries:

rng ... range queries ... $0 \leqslant r_q \leqslant r$, $0 \leqslant b_q \leqslant b+b_{ov}$

- conjunction of inequalities, on one or more attributes, e.g.
- **`select * from R where age≥18 and age≤21`** (1-d)
- **`select * from R where 18≤age≤21 and 160≤height≤190`** (n-d)

pat ... pattern-based queries ... $0 \leqslant r_q \leqslant r$, $0 \leqslant b_q \leqslant b+b_{ov}$

- string-based matching using **`like`** or regular expressions
- **`select * from R where name like '%oo%'`**
- **`select * from R where name ~ '^Smi'`**

<< ^ >>

❖ Varieties of Selection (cont)

More categories of selection queries:

sim ... similarity matching ... in theory, $r_q = r$ ... everything matches to some degree

- uses "similarity" measure (0 ⩽ *sim* ⩽ 1, 0=different, 1=identical)
- **`select * from`** *Images* **`where similar to`** *SampleImage*
- results are ranked by *sim* value, from most to least similar
- can become a filter via
  - threshold ... only items where *sim* ⩾ min similarity
  - top-k ... *k* items with highest similarities

We focus on one, pmr and rng queries, but will discuss others

<< ^

❖ Implementing Select Efficiently

Two basic approaches:

- physical arrangement of tuples
  - sorting (search strategy)
  - hashing (static, dynamic, *n*-dimensional)
- additional indexing information
  - index files (primary, secondary, trees)
  - signatures (superimposed, disjoint)

Our analysis assumes 1 input buffer available for each relation.

If more buffers are available, most methods benefit.

Produced: 6 Mar 2021