

# Sorted Files

---


- Sorted Files
- Selection in Sorted Files
- Insertion into Sorted Files
- Deletion from Sorted Files

## ❖ Sorted Files


Records stored in file in order of some field **k** (the sort key).

Makes searching more efficient; makes insertion less efficient

E.g. assume  $c = 4$




*insert g*



## ❖ Sorted Files (cont)

In order to mitigate insertion costs, use overflow pages.



Total number of overflow pages =  $b_{ov}$ .

Average overflow chain length =  $O_v = b_{ov} / b$ .

**Bucket** = data page + its overflow page(s)

## ❖ Selection in Sorted Files

For *one* queries on sort key, use binary search.

```
// select * from R where k = val  (sorted on R.k)
lo = 0; hi = nPages(rel)-1
while (lo <= hi) {
    mid = (lo+hi) / 2;  // int division with truncation
    (tup, loVal, hiVal) = searchBucket(rel, mid, x, val);
    if (tup != NULL) return tup;
    else if (val < loVal) hi = mid - 1;
    else if (val > hiVal) lo = mid + 1;
    else return NOT_FOUND;
}
return NOT_FOUND;
```

where **rel** is relation handle, **mid**, **lo**, **hi** are page indexes,  
**k** is a field/attr, **val**, **loVal**, **hiVal** are values for **k**

## ❖ Selection in Sorted Files (cont)

Search a page and its overflow chain for a key value

```
searchBucket(rel,p,k,val)
{
    get_page(rel,p,buf);
    (tup,min,max) = searchPage(buf,k,val,+INF,-INF)
    if (tup != NULL) return(tup,min,max);
    ovf = openOvFile(f);
    ovp = overflow(buf);
    while (tup == NULL && ovp != NO_PAGE) {
        get_page(ovf,ovp,buf);
        (tup,min,max) = searchPage(buf,k,val,min,max)
        ovp = overflow(buf);
    }
    return (tup,min,max);
}
```

Assumes each page contains index of next page in Ov chain

## ❖ Selection in Sorted Files (cont)

Search within a page for key; also find min/max key values

```
searchPage(buf, k, val, min, max)
{
    res = NULL;
    for (i = 0; i < nTuples(buf); i++) {
        tup = get_tuple(buf, i);
        if (tup.k == val) res = tup;
        if (tup.k < min) min = tup.k;
        if (tup.k > max) max = tup.k;
    }
    return (res, min, max);
}
```

## ❖ Selection in Sorted Files (cont)

The above method treats each bucket like a single large page.

Cases:

- best: find tuple in first data page we read
- worst: full binary search, and not found
  - examine  $\log_2 b$  data pages
  - plus examine all of their overflow pages
- average: examine some data pages + their overflow pages

$Cost_{one}$ : Best = 1 Worst =  $\log_2 b + b_{ov}$


Average case cost analysis needs assumptions (e.g. data distribution)

## ❖ Selection in Sorted Files (cont)

For *pmr* query, on non-unique attribute  $k$ , where file is sorted on  $k$

- tuples containing  $k$  may span several pages

E.g. **select \* from R where  $k = 2$**



Begin by locating a page  $p$  containing  $k=val$  (as for *one* query).

Scan backwards and forwards from  $p$  to find matches.

Thus,  $Cost_{pmr} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$




## ❖ Selection in Sorted Files (cont)

For *range* queries on unique sort key (e.g. primary key):

- use binary search to find lower bound
- read sequentially until reach upper bound

E.g. **select \* from R where k >= 5 and k <= 13**




$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

## ❖ Selection in Sorted Files (cont)

For *range* queries on non-unique sort key, similar method to *pmr*.

- binary search to find lower bound
- then go backwards to start of run
- then go forwards to last occurrence of upper-bound

E.g. **select \* from R where k >= 2 and k <= 6**



$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

## ❖ Selection in Sorted Files (cont)

So far, have assumed query condition involves sort key  $k$ .

But what about **select \* from R where j = 100.0?**

If condition contains attribute  $j$ , not the sort key

- file is unlikely to be sorted by  $j$  as well
- sortedness gives no searching benefits

$Cost_{one}$ ,  $Cost_{range}$ ,  $Cost_{pmr}$  as for heap files

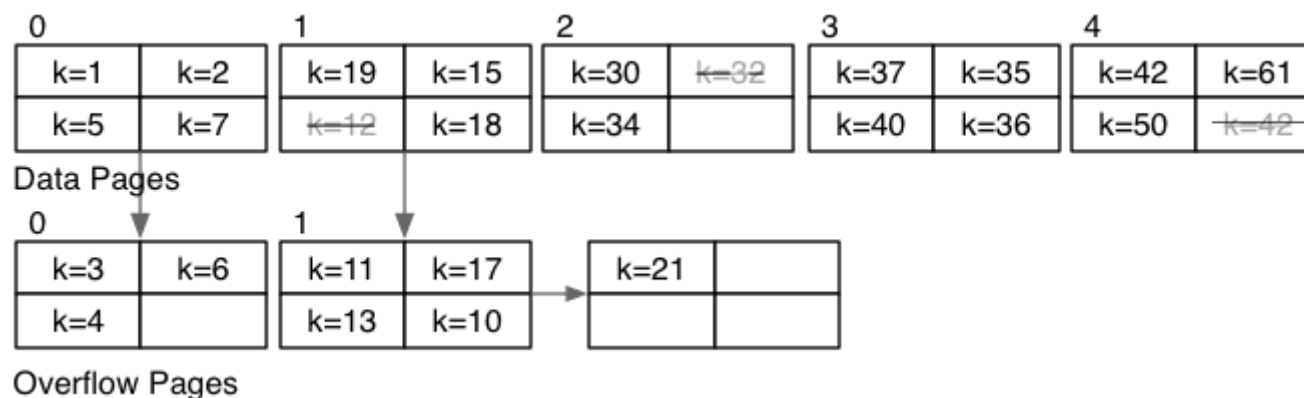
## ❖ Insertion into Sorted Files

Insertion approach:

- find appropriate page for tuple (via binary search)
- if page not full, insert into page
- otherwise, insert into next overflow page with space

Thus,  $Cost_{insert} = Cost_{one} + \delta_w$  (where  $\delta_w = 1$  or  $2$ )

Consider insertions of  $k=33$ ,  $k=25$ ,  $k=99$  into:



## ❖ Deletion from Sorted Files

E.g. **delete from R where k = 2**

**Deletion strategy:**

- find matching tuple(s)
- mark them as deleted

Cost depends on **selectivity** of selection condition

Recall: selectivity determines  $b_q$  (# pages with matches)

Thus,  $Cost_{delete} = Cost_{select} + b_{qw}$

Produced: 7 Mar 2021