

Week 4 Exercises

- Things to Note
- Assignment 1
- Exercise: File Merging
- Exercise: Sort-based Projection
- Exercise: Hash-based Projection
- Exercise: Query Types
- Exercise: Cost of Deletion in Heaps
- Exercise: Searching in Sorted File
- Exercise: Optimising Sorted-file Search
- Exercise: Insertion into Static Hashed File
- Exercise: Bit Manipulation
- Exercise: Insertion into Linear Hashed File

❖ Things to Note

- "Disk quota exceeded"? Check your disk usage on /srvr/
 - use the **du** command to find "hot spots"
- This weekend is Census Weekend ...
 - if you want to drop a course and not pay for it, drop it now
 - if you want to drop COMP9315, drop it in MyUNSW and email me

❖ Assignment 1

- three data usages: readable, computable, storeable
- *could* use the same representation for all three
- or, more likely
 - readable → storeable in **intSet_in()**
 - storeable → readable in **intSet_out()**
 - storeable → computable in e.g. **intSet_member()**
 - *could* use same representation for storeable and computable
- reminder: complex in-memory structures (e.g. BSTs) are not storeable
 - need to serialize/flatten them to make a storeable byte-sequence

❖ Exercise: File Merging

Implement a merging algorithm

- for two sorted files, using 3 buffers, with $b_1=5$, $b_2=3$
- for one unsorted file, using 3 buffers, with $b = 12$
- for one unsorted file, using 5 buffers, with $b = 27$

Assume that we have functions

- **get_page(rel, pid, buf)** ... read specified page into buffer
- **put_page(rel, pid, buf)** ... write a page to disk, at position pid
- **clear_page(rel, buf)** ... make page have zero tuples
- **sort_page(buf)** ... in-memory sort of tuples in page
- **nPages(rel), nTuples(buf), get_tuple(buf, tid)**

❖ Exercise: Sort-based Projection

Consider a table $R(x,y,z)$ with tuples:

Page 0:	(1,1,'a')	(11,2,'a')	(3,3,'c')
Page 1:	(13,5,'c')	(2,6,'b')	(9,4,'a')
Page 2:	(6,2,'a')	(17,7,'a')	(7,3,'b')
Page 3:	(14,6,'a')	(8,4,'c')	(5,2,'b')
Page 4:	(10,1,'b')	(15,5,'b')	(12,6,'b')
Page 5:	(4,2,'a')	(16,9,'c')	(18,8,'c')

SQL: **create T as (select distinct y from R)**

Assuming:

- 3 memory buffers, 2 for input, one for output
- pages/buffers hold 3 **R** tuples (i.e. $c_R=3$), 6 **T** tuples (i.e. $c_T=6$)

Show how sort-based projection would execute this statement.

❖ Exercise: Hash-based Projection

Consider a table $R(x,y,z)$ with tuples:

```

Page 0:  (1,1,'a')    (11,2,'a')    (3,3,'c')
Page 1:  (13,5,'c')   (2,6,'b')    (9,4,'a')
Page 2:  (6,2,'a')    (17,7,'a')   (7,3,'b')
Page 3:  (14,6,'a')   (8,4,'c')    (5,2,'b')
Page 4:  (10,1,'b')   (15,5,'b')   (12,6,'b')
Page 5:  (4,2,'a')    (16,9,'c')   (18,8,'c')
-- and then the same tuples repeated for pages 6-11

```

SQL: **create T as (select distinct y from R)**

Assuming:

- 4 memory buffers, one for input, 3 for partitioning
- pages/buffers hold 3 **R** tuples (i.e. $c_R=3$), 4 **T** tuples (i.e. $c_T=4$)
- hash functions: $h1(x) = x\%3$, $h2(x) = (x\%4)\%3$

Show how hash-based projection would execute this statement.

❖ Exercise: Query Types

Using the relation:

```
create table Courses (  
    id          integer primary key,  
    code        char(8),    -- e.g. 'COMP9315'  
    title       text,       -- e.g. 'Computing 1'  
    year        integer,    -- e.g. 2000..2016  
    convenor    integer references Staff(id),  
    constraint once_per_year unique (code,year)  
);
```

give examples of each of the following query types:

1. a 1-d *one* query, an n-d *one* query
2. a 1-d *pmr* query, an n-d *pmr* query
3. a 1-d *range* query, an n-d *range* query

Suggest how many solutions each might produce ...

❖ Exercise: Cost of Deletion in Heaps

Consider the following queries ...

```
delete from Employees where id = 12345    -- one  
delete from Employees where dept = 'Marketing'  -- pmr  
delete from Employees where 40 <= age and age < 50  -- range
```

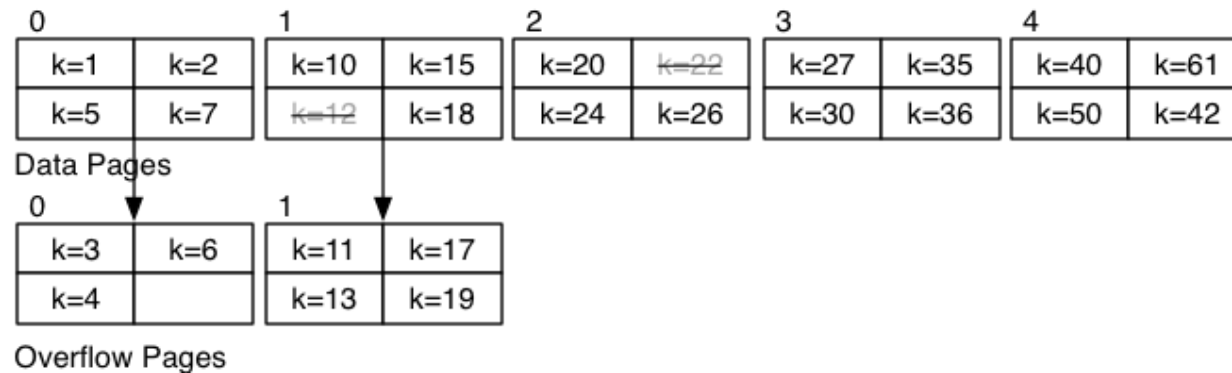
Show how each will be executed and estimate the cost, assuming:

- $b = 100$, $b_{q2} = 3$, $b_{q3} = 20$

State any other assumptions.

❖ Exercise: Searching in Sorted File

Consider this sorted file with overflows ($b=5, c=4$):



Compute the cost for answering each of the following:

- **select * from R where k = 24**
- **select * from R where k = 3**
- **select * from R where k = 14**
- **select max(k) from R**

❖ Exercise: Optimising Sorted-file Search

The **searchBucket (f, p, k, val)** function requires:

- read the p^{th} page from data file
- scan it to find a match and min/max k values in page
- while no match, repeat the above for each overflow page
- if we find a match in any page, return it
- otherwise, remember min/max over all pages in bucket

Suggest an optimisation that would improve **searchBucket ()** performance for most buckets.

❖ Exercise: Insertion into Static Hashed File

Consider a file with $b=4$, $c=3$, $d=2$, $h(x) = \text{bits}(d, \text{hash}(x))$

Insert tuples in alpha order with the following keys and hashes:

k	$\text{hash}(k)$	k	$\text{hash}(k)$	k	$\text{hash}(k)$	k	$\text{hash}(k)$
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

❖ Exercise: Bit Manipulation

1. Write a function to display **uint32** values as **01010110...**

```
char *showBits(uint32 val, char *buf);
```

Analogous to **gets()** (assumes supplied buffer is large enough)

2. Write a function to extract the d bits of a **uint32**

```
uint32 bits(int d, uint32 val);
```

If $d > 0$, gives low-order bits; if $d < 0$, gives high-order bits

❖ Exercise: Insertion into Linear Hashed File

Consider a file with $b=4$, $c=3$, $d=2$, $sp=0$, $hash(x)$ as below

Insert tuples in alpha order with the following keys and hashes:

k	$hash(k)$	k	$hash(k)$	k	$hash(k)$	k	$hash(k)$
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

Produced: 11 Mar 2021