

Nested-loop Join

- Join Example
- Nested Loop Join
- Block Nested Loop Join
- Cost on Example Query
- Block Nested Loop Join
- Index Nested Loop Join

❖ Join Example

SQL query on student/enrolment database:

```
select E.subj, S.name
from   Student S join Enrolled E on (S.id = E.stude)
order by E.subj
```

And its relational algebra equivalent:

$Sort[subj] (Project[subj,name] (Join[id=stude](Student, Enrolled)))$

Database: $r_S = 20000$, $c_S = 20$, $b_S = 1000$, $r_E = 80000$, $c_E = 40$, $b_E = 2000$

We are interested only in the cost of *Join*, with N buffers

❖ Nested Loop Join

Basic strategy ($R.a \bowtie S.b$):

```
Result = {}
for each page i in R {
  pageR = getPage(R,i)
  for each page j in S {
    pageS = getPage(S,j)
    for each pair of tuples  $t_R, t_S$ 
      from pageR, pageS {
        if ( $t_R.a == t_S.b$ )
          Result = Result  $\cup$  ( $t_R:t_S$ )
      }
    }
  }
```

Needs input buffers for R and S, output buffer for "joined" tuples

Terminology: R is outer relation, S is inner relation


Cost = $b_R \cdot b_S$... ouch!

COMP9315 21T1 ◇ Nested-loop Join ◇ [2/8]

❖ Block Nested Loop Join

Method (for N memory buffers):

- read $N-2$ -page chunk of R into memory buffers
- for each S page
 check join condition on all (t_R, t_S) pairs in buffers
- repeat for all $N-2$ -page chunks of R



❖ Block Nested Loop Join (cont)

Best-case scenario: $b_R \leq N-2$

- read b_R pages of relation R into buffers
- while whole R is buffered, read b_S pages of S

$$\text{Cost} = b_R + b_S$$

Typical-case scenario: $b_R > N-2$

- read $\text{ceil}(b_R/(N-2))$ chunks of pages from R
- for each chunk, read b_S pages of S

$$\text{Cost} = b_R + b_S \cdot \text{ceil}(b_R/N-2)$$

Note: always requires $r_R \cdot r_S$ checks of the join condition

❖ Cost on Example Query

With $N = 12$ buffers, and S as outer and E as inner

- $\text{Cost} = b_S + b_E \cdot \text{ceil}(b_S / (N - 2)) = 1000 + 2000 \cdot \text{ceil}(1000 / 10) = 201000$

With $N = 12$ buffers, and E as outer and S as inner

- $\text{Cost} = b_E + b_S \cdot \text{ceil}(b_E / (N - 2)) = 2000 + 1000 \cdot \text{ceil}(2000 / 10) = 202000$

With $N = 102$ buffers, and S as outer and E as inner

- $\text{Cost} = b_S + b_E \cdot \text{ceil}(b_S / (N - 2)) = 1000 + 2000 \cdot \text{ceil}(1000 / 100) = 21000$

With $N = 102$ buffers, and E as outer and S as inner

- $\text{Cost} = b_E + b_S \cdot \text{ceil}(b_E / (N - 2)) = 2000 + 1000 \cdot \text{ceil}(2000 / 100) = 22000$

COMP9315 21T1 ◇ Nested-loop Join ◇ [5/8]

❖ Block Nested Loop Join

Why block nested loop join is actually useful in practice ...

Many queries have the form

```
select *  
from   R join S on (R.i = S.j)  
where  R.x = K
```

This would typically be evaluated as

```
Tmp = Sel[x=K](R)  
Res = Join[i=j](Tmp, S)
```

If **Tmp** is small \Rightarrow may fit in memory (in small #buffers)

❖ Index Nested Loop Join

A problem with nested-loop join:

- needs repeated scans of *entire* inner relation S

If there is an index on S , we can avoid such repeated scanning.

Consider $Join_{[i=j]}(R, S)$:

```
for each tuple r in relation R {  
    use index to select tuples  
        from S where  $s.j = r.i$   
    for each selected tuple s from S {  
        add (r,s) to result  
    }  
}
```

❖ Index Nested Loop Join (cont)

This method requires:

- one scan of R relation (b_R)
 - only one buffer needed, since we use R tuple-at-a-time
- for each tuple in R (r_R), one index lookup on S
 - cost depends on type of index and number of results
 - best case is when each $R.i$ matches few S tuples

Cost = $b_R + r_R \cdot Sel_S$ (Sel_S is the cost of performing a select on S).

Typical $Sel_S = 1-2$ (hashing) .. b_q (unclustered index)

Trade-off: $r_R \cdot Sel_S$ vs $b_R \cdot b_S$, where $b_R \ll r_R$ and $Sel_S \ll b_S$

Produced: 28 Mar 2021