

Hash Join

- Hash Join
- Simple Hash Join
- Grace Hash Join
- Hybrid Hash Join
- Costs for Join Example
- Join in PostgreSQL

❖ Hash Join

Basic idea:

- use hashing as a technique to partition relations
- to avoid having to consider all pairs of tuples

Requires sufficient memory buffers

- to hold substantial portions of partitions
- (preferably) to hold largest partition of outer relation

Other issues:

- works only for equijoin **$R.i = S.j$** (but this is a common case)
- susceptible to data skew (or poor hash function)

Variations: *simple, grace, hybrid.*

❖ Simple Hash Join

Basic approach:

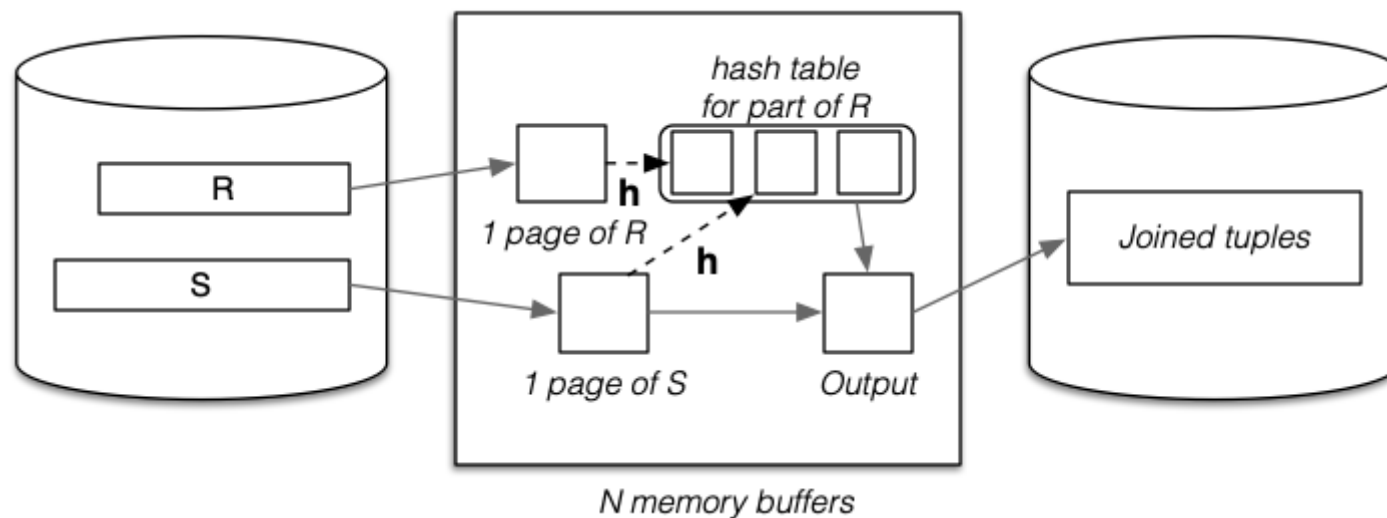
- hash part of outer relation R into memory buffers (build)
- scan inner relation S , using hash to search (probe)
 - if $R.i=S.j$, then $h(R.i)=h(S.j)$ (hash to same buffer)
 - only need to check one memory buffer for each S tuple
- repeat until whole of R has been processed

No overflows allowed in in-memory hash table

- works best with uniform hash function
- can be adversely affected by data/hash skew

❖ Simple Hash Join (cont)

Data flow in hash join:



❖ Simple Hash Join (cont)

Algorithm for simple hash join $Join[R.i=S.j](R,S)$:

```

for each tuple r in relation R {
  if (buffer[h(R.i)] is full) {
    for each tuple s in relation S {
      for each tuple rr in buffer[h(S.j)] {
        if ((rr,s) satisfies join condition) {
          add (rr,s) to result
        }
      }
    }
    clear all hash table buffers
  }
  insert r into buffer[h(R.i)]
}

```

Best case: #join tests $\leq r_S \cdot c_R$ (cf. nested-loop $r_S \cdot r_R$)

❖ Simple Hash Join (cont)

Cost for simple hash join ...

Best case: all tuples of R fit in the hash table

- Cost = $b_R + b_S$
- Same page reads as block nested loop, but less join tests

Good case: refill hash table m times (where $m \geq \text{ceil}(b_R / (N-3))$)

- Cost = $b_R + m.b_S$
- More page reads than block nested loop, but less join tests

Worst case: everything hashes to same page

- Cost = $b_R + b_R.b_S$

❖ Grace Hash Join

Basic approach (for $R \bowtie S$):

- partition both relations on join attribute using hashing ($h1$)
- load each partition of R into $N-3$ *buffer hash table ($h2$)
- scan through corresponding partition of S to form results
- repeat until all partitions exhausted

For best-case cost ($O(b_R + b_S)$):


- need $\geq \sqrt{b_R}$ buffers to hold largest partition of outer relation

If $< \sqrt{b_R}$ buffers or poor hash distribution

- need to scan some partitions of S multiple times


❖ Grace Hash Join (cont)

Partition phase (applied to both R and S):



❖ Grace Hash Join (cont)

Probe/join phase:



The second hash function (**h2**) simply speeds up the matching process. Without it, would need to scan entire *R* partition for each record in *S* partition.

❖ Grace Hash Join (cont)

Cost of grace hash join:

- #pages in all partition files of $Rel \approx b_{Rel}$ (maybe slightly more)
- partition relation R ... Cost = $read(b_R) + write(\approx b_R) = 2b_R$
- partition relation S ... Cost = $read(b_S) + write(\approx b_S) = 2b_S$
- probe/join requires one scan of each (partitioned) relation
Cost = $b_R + b_S$
- all hashing and comparison occurs in memory \Rightarrow tiny cost

$$\text{Total Cost} = 2b_R + 2b_S + b_R + b_S = 3(b_R + b_S)$$

❖ Hybrid Hash Join

A variant of grace hash join if we have $\sqrt{b_R} < N < b_R + 2$

- create $k \ll N$ partitions, 1 in memory, $k-1$ on disk
- buffers: 1 input, $k-1$ output, $p = N - k - 2$ for in-memory partition

When we come to scan and partition S relation

- any tuple with hash 0 can be resolved (using in-memory partition)
- other tuples are written to one of k partition files for S


Final phase is same as grace join, but with only $k-1$ partitions.

Comparison:

- grace hash join creates $N-1$ partitions on disk
- hybrid hash join creates 1 (memory) + $k-1$ (disk) partitions


❖ Hybrid Hash Join (cont)

First phase of hybrid hash join (partitioning R):



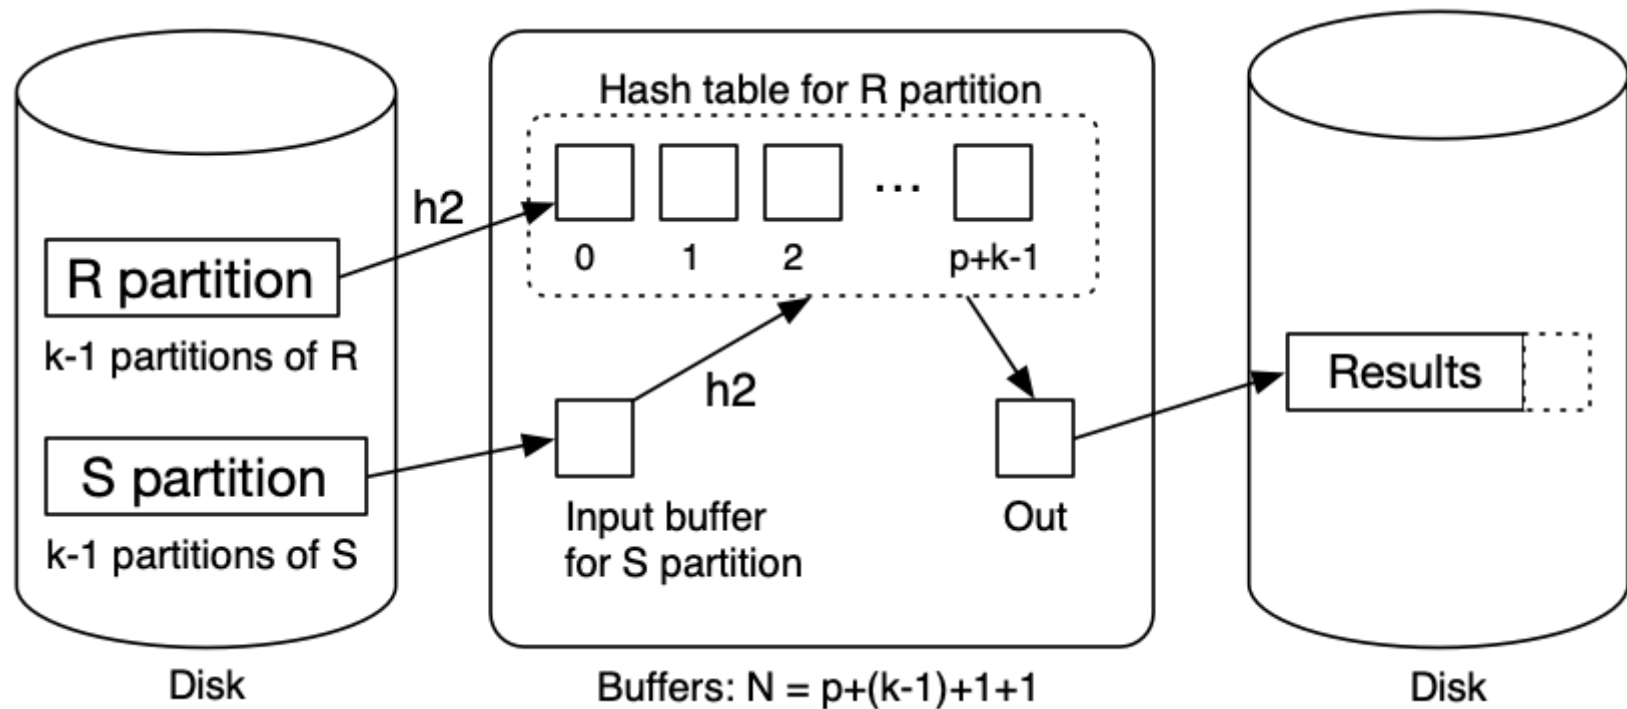
❖ Hybrid Hash Join (cont)

Next phase of hybrid hash join (partitioning S):



❖ Hybrid Hash Join (cont)

Final phase of hybrid hash join (finishing join):



❖ Hybrid Hash Join (cont)

Some observations:

- with k partitions, each partition has expected size $\text{ceil}(b_R/k)$
- holding 1 partition in memory needs $\text{ceil}(b_R/k)$ buffers
- trade-off between in-memory partition space and #partitions

Other notes:

- if $N = b_R + 2$, using block nested loop join is simpler
- cost depends on N (but less than grace hash join)

For k partitions, Cost = $(3 - 2/k) \cdot (b_R + b_S)$

❖ Costs for Join Example

SQL query on student/enrolment database:

```
select E.subj, S.name
from   Student S join Enrolled E on (S.id = E.stude)
order by E.subj
```

And its relational algebra equivalent:

$Sort[subj] (Project[subj,name] (Join[id=stude](Student, Enrolled)))$

Database: $r_S = 20000$, $c_S = 20$, $b_S = 1000$, $r_E = 80000$, $c_E = 40$, $b_E = 2000$

We are interested only in the cost of *Join*, with N buffers

❖ Costs for Join Example (cont)

Costs for hash join variants on example ($N=103$):

Hash Join Method	Cost Analysis	Cost
Hybrid Hash Join	$(3-2/k).(b_S+b_E) = 2.8((1000+2000)$ assuming $k = 10$... and one partition fits in 91 pages	8700
Grace Hash Join	$3(b_S+b_E) = 3(1000+2000)$	9000
Simple Hash Join	$b_S + b_E.\text{ceil}(b_R/(N-3)) =$ $1000 + \text{ceil}(1000/100).2000 = 1000 + 10.2000$	21000
Sort-merge Join	$\text{sort}(S) + \text{sort}(E) + b_S + b_E =$ $2.1000.2 + 2.2000.2 + 1000 + 2000$	11000
Nested-loop Join	$b_S + b_E.\text{ceil}(b_S/(N-2)) =$ $1000 + 2000.\text{ceil}(1000/101) = 1000 + 10.2000$	21000

COMP9315 21T1 ◇ Hash Join ◇ [16/17]

❖ Join in PostgreSQL

Join implementations are under: **src/backend/executor**

PostgreSQL supports three kinds of join:

- nested loop join (**nodeNestloop.c**)
- sort-merge join (**nodeMergejoin.c**)
- hash join (**nodeHashjoin.c**) (hybrid hash join)

Query optimiser chooses appropriate join, by considering

- physical characteristics of tables being joined
- estimated selectivity (likely number of result tuples)

Produced: 25 Apr 2021