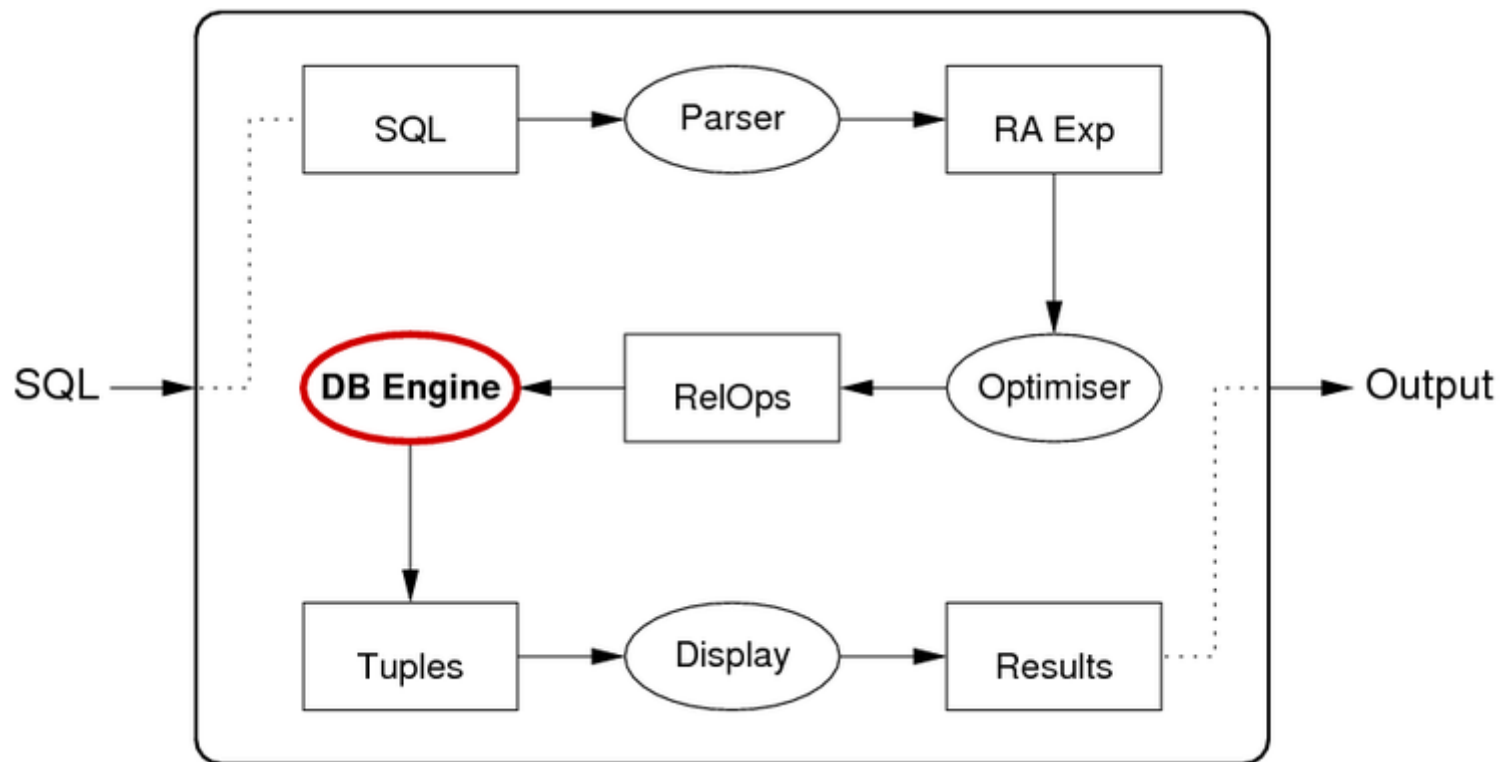


Query Execution

- Query Execution
- Materialization
- Pipelining
- Iterators (reminder)
- Pipelining Example
- Disk Accesses
- PostgreSQL Query Execution
- PostgreSQL Executor
- Example PostgreSQL Execution

❖ Query Execution

Query execution: applies evaluation plan → result tuples



❖ Query Execution (cont)

Example of query translation:

```
select s.name, s.id, e.course, e.mark
from   Student s, Enrolment e
where  e.student = s.id and e.semester = '05s2';
```

maps to

$$\pi_{name,id,course,mark}(Stu \bowtie_{e.student=s.id} (\sigma_{semester=05s2} Enr))$$

maps to

```
Temp1  = BtreeSelect[semester=05s2](Enr)
Temp2  = HashJoin[e.student=s.id](Stu,Temp1)
Result = Project[name,id,course,mark](Temp2)
```

❖ Query Execution (cont)

A query execution plan:

- consists of a **collection of RelOps**
- executing together to produce a set of result tuples

Results may be passed from one operator to the next:

- **materialization** ... writing results to disk and reading them back
- **pipelining** ... generating and passing via memory buffers

❖ Materialization

Steps in **materialization** between two operators

- first operator reads input(s) and writes results to disk
- next operator treats tuples on disk as its input
- in essence, the **Temp** tables are produced as real tables

Advantage:

- intermediate results can be placed in a file structure
(which can be chosen to speed up execution of subsequent operators)

Disadvantage:

- requires disk space/writes for intermediate results
- requires disk access to read intermediate results

❖ Pipelining

How **pipelining** is organised between two operators:

- operators execute "concurrently" as producer/consumer pairs
- structured as interacting iterators (open; while(next); close)

Advantage:

- no requirement for disk access (results passed via memory buffers)

Disadvantage:

- higher-level operators access inputs via linear scan, or
- requires sufficient memory buffers to hold all outputs

❖ Iterators (reminder)

Iterators provide a "stream" of results:

- **`iter = startScan(params)`**
 - set up data structures for iterator (create state, open files, ...)
 - *params* are specific to operator (e.g. reln, condition, #buffers, ...)
- **`tuple = nextTuple(iter)`**
 - get the next tuple in the iteration; return null if no more
- **`endScan(iter)`**
 - clean up data structures for iterator


Other possible operations: reset to specific point, restart, ...

❖ Pipelining Example

Consider the query:

```
select s.id, e.course, e.mark
from   Student s, Enrolment e
where  e.student = s.id and
       e.semester = '05s2' and s.name = 'John';
```

Evaluated via communication between RA tree nodes:



Note: likely that projection is combined with join in PostgreSQL

COMP9315 21T1 ◇ Query Execution ◇ [7/14]

❖ Disk Accesses

Pipelining cannot avoid all disk accesses.

Some operations use multiple passes (e.g. merge-sort, hash-join).

- data is written by one pass, read by subsequent passes

Thus ...

- **within** an operation, disk reads/writes are possible
- **between** operations, no disk reads/writes are needed

❖ PostgreSQL Query Execution

Defs: **src/include/executor** and **src/include/nodes**

Code: **src/backend/executor**

PostgreSQL uses pipelining (as much as possible) ...

- query plan is a tree of **Plan** nodes
- each type of node implements one kind of RA operation
(node implements specific access method via iterator interface)
- node types e.g. **Scan**, **Group**, **Indexscan**, **Sort**, **HashJoin**
- execution is managed via a tree of **PlanState** nodes
(mirrors the structure of the tree of Plan nodes; holds execution state)

❖ PostgreSQL Executor

Modules in **src/backend/executor** fall into two groups:

execXXX (e.g. execMain, execProcnode, execScan)

- implement generic control of plan evaluation (execution)
- provide overall plan execution and dispatch to node iterators

nodeXXX (e.g. nodeSeqscan, nodeNestloop, nodeGroup)


- implement iterators for specific types of RA operators
- typically contains **ExecInitXXX**, **ExecXXX**, **ExecEndXXX**

❖ Example PostgreSQL Execution

Consider the query:

```
-- get manager's age and # employees in Shoe department  
select e.age, d.nemps  
from   Departments d, Employees e  
where  e.name = d.manager and d.name = 'Shoe'
```

and its execution plan tree



❖ Example PostgreSQL Execution (cont)

Initially **InitPlan()** invokes **ExecInitNode()** on plan tree root.

ExecInitNode() sees a **NestedLoop** node ...

so dispatches to **ExecInitNestLoop()** to set up iterator

then invokes **ExecInitNode()** on left and right sub-plans

in left subPlan, **ExecInitNode()** sees an **IndexScan** node

so dispatches to **ExecInitIndexScan()** to set up iterator

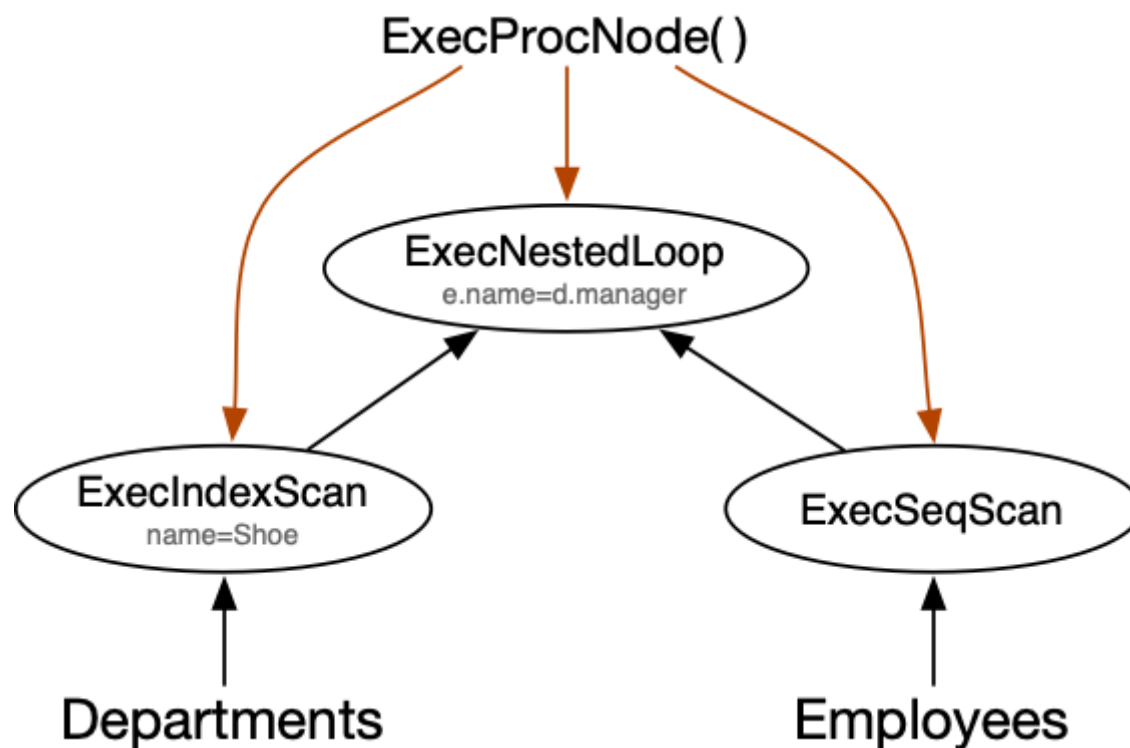
in right sub-plan, **ExecInitNode()** sees a **SeqScan** node

so dispatches to **ExecInitSeqScan()** to set up iterator

Result: a plan *state* tree with same structure as plan tree.

❖ Example PostgreSQL Execution (cont)

Plan state tree (collection of iterators):



❖ Example PostgreSQL Execution (cont)

Then **ExecutePlan()** repeatedly invokes **ExecProcNode()**.

ExecProcNode() sees a **NestedLoop** node ...

so dispatches to **ExecNestedLoop()** to get next tuple

which invokes **ExecProcNode()** on its sub-plans

in left sub-plan, **ExecProcNode()** sees an **IndexScan** node

so dispatches to **ExecIndexScan()** to get next tuple

if no more tuples, return END

for this tuple, invoke **ExecProcNode()** on right sub-plan

ExecProcNode() sees a **SeqScan** node

so dispatches to **ExecSeqScan()** to get next tuple

check for match and return joined tuples if found

continue scan until end

reset right sub-plan iterator

Result: stream of result tuples returned via **ExecutePlan()**

Produced: 6 Apr 2021