

Lock-based Concurrency Control

- Lock-based Concurrency Control
- Two-Phase Locking
- Problems with Locking
- Deadlock

❖ Lock-based Concurrency Control

Requires read/write **lock** operations which act on database objects.

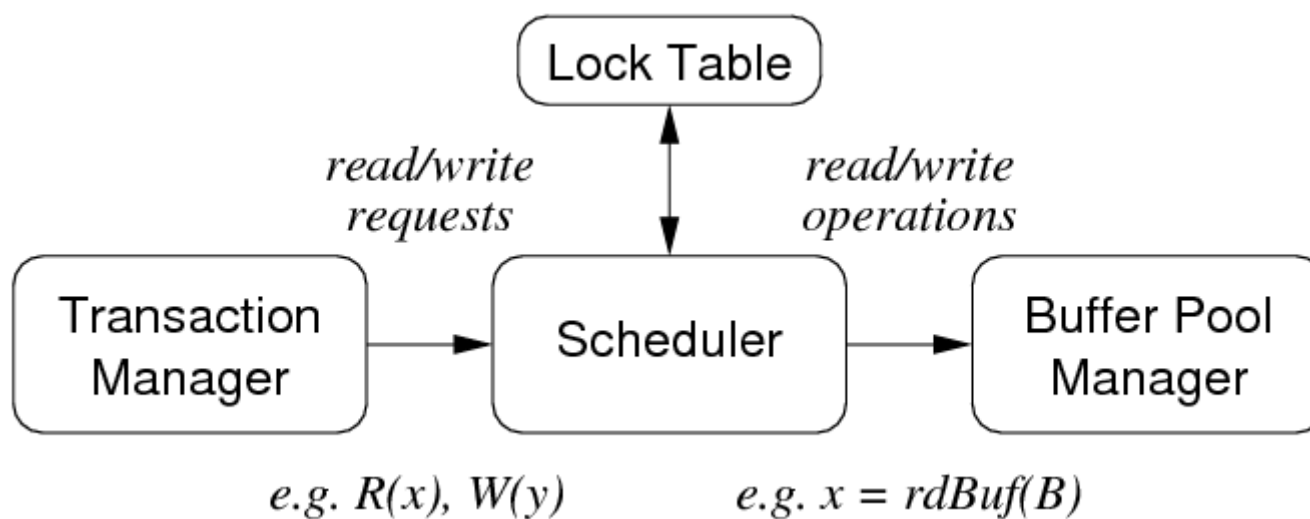
Synchronise access to shared DB objects via these rules:

- before reading X , get read (shared) lock on X
- before writing X , get write (exclusive) lock on X
- a tx attempting to get a read lock on X is blocked if another tx already has write lock on X
- a tx attempting to get a write lock on X is blocked if another tx has any kind of lock on X

These rules alone do not guarantee serializability.

❖ Lock-based Concurrency Control (cont)

Locks introduce additional mechanisms in DBMS:



The Lock Manager manages the locks requested by the scheduler

❖ Lock-based Concurrency Control (cont)

Lock table entries contain:

- object being locked (DB, table, tuple, field)
- type of lock: read/shared, write/exclusive
- FIFO queue of tx's requesting this lock
- count of tx's currently holding lock (max 1 for write locks)

Lock and unlock operations *must* be atomic.

Lock **upgrade**:

- if a tx holds a read lock, and it is the only tx holding that lock
- then the lock can be converted into a write lock

❖ Lock-based Concurrency Control (cont)

Consider the following schedule, using locks:

T1(a) : $L_r(Y)$ R(Y) continued

T2(a) : $L_r(X)$ R(X) U(X) continued

T1(b) : U(Y) $L_w(X)$ W(X) U(X)

T2(b) : $L_w(Y) \dots W(Y)$ U(Y)

(where L_r = read-lock, L_w = write-lock, U = unlock)

Locks correctly ensure controlled access to **x** and **y**.

Despite this, the schedule is not serializable. (Ex: prove this)

❖ Two-Phase Locking

To guarantee serializability, we require an additional constraint:

- in every tx, all *lock* requests precede all *unlock* requests

Each transaction is then structured as:

- **growing** phase where locks are acquired
- **action** phase where "real work" is done
- **shrinking** phase where locks are released

Clearly, this reduces potential concurrency ...

❖ Problems with Locking

Appropriate locking can guarantee serializability..

However, it also introduces potential undesirable effects:

- **Deadlock**
 - No transactions can proceed; each waiting on lock held by another.
- **Starvation**
 - One transaction is permanently "frozen out" of access to data.
- **Reduced performance**
 - Locking introduces delays while waiting for locks to be released.

❖ Deadlock

Deadlock occurs when two tx's wait for a lock on an item held by the other.

Example:

T1:	$L_W(A)$	$R(A)$		$L_W(B)$
T2:			$L_W(B)$	$R(B)$	$L_W(A)$

How to deal with deadlock?

- prevent it happening in the first place
- let it happen, detect it, recover from it

❖ Deadlock (cont)

Handling deadlock involves forcing a transaction to "back off"

- select process to roll back
 - choose on basis of how far tx has progressed, # locks held, ...
- roll back the selected process
 - how far does this it need to be rolled back?
 - worst-case scenario: abort one transaction, then retry
- prevent starvation
 - need methods to ensure that same tx isn't always chosen

❖ Deadlock (cont)

Methods for managing deadlock

- **timeout** : set max time limit for each tx
- **waits-for graph** : records T_j waiting on lock held by T_k
 - *prevent* deadlock by checking for new cycle \Rightarrow abort T_i
 - *detect* deadlock by periodic check for cycles \Rightarrow abort T_i
- **timestamps** : use tx start times as basis for priority
 - scenario: T_j tries to get lock held by T_k ...
 - **wait-die**: if $T_j < T_k$, then T_j waits, else T_j rolls back
 - **wound-wait**: if $T_j < T_k$, then T_k rolls back, else T_j waits

❖ Deadlock (cont)

Properties of deadlock handling methods:

- both wait-die and wound-wait are fair
- wait-die tends to
 - roll back tx's that have done little work
 - but rolls back tx's more often
- wound-wait tends to
 - roll back tx's that may have done significant work
 - but rolls back tx's less often
- timestamps easier to implement than waits-for graph
- waits-for minimises roll backs because of deadlock

Produced: 14 Apr 2021