

# Quiz 5

Deadline	Friday, 24 April 2020 at 11:59PM
Latest Submission	Friday, 17 April 2020 at 5:46PM
Maximum Mark	4

## Question 1 (1 mark)

Consider a table of machine parts (*Parts(id,name,colour,size,cost)*), about which we have the following information:

```
db=# select count(distinct size) from Parts;  
count
```

-----  
5

```
db=# select min(size) from Parts;  
min
```

-----  
tiny

```
db=# select max(size) from Parts;  
max
```

-----  
huge

```
db=# select count(*) from Parts;  
count
```

-----  
6507

Approximately how many tuples would you expect to find in the result of the following query?

```
select * from Parts where size = 'huge';
```

Assume a uniform distribution of part sizes.

(a) <input type="radio"/>	650
(b) <input type="radio"/>	3253
(c) <input type="radio"/>	6506
(d) <input type="radio"/>	None of the other options is correct.

(e) <input checked="" type="radio"/>	1301
--------------------------------------	------

### Question 2 (1 mark)

Consider the three relations  $R(id,a,b)$ ,  $S(rid,tid,c)$ ,  $T(id,d,e,f)$  where the  $id$  attributes are primary keys, and  $S.rid$  is a foreign key referencing  $R.id$  and  $S.tid$  is a foreign key referencing  $T.id$ .

Consider also the following join on these three relations:

```
select a,b,c,d,e,f from R, S, T where R.id = S.rid and T.id = S.tid
```

Which of the following does **not** represent a possible join order for the above SQL statement?

(a) <input type="radio"/>	$(R \bowtie S) \bowtie T$
(b) <input type="radio"/>	$R \bowtie (S \bowtie T)$
(c) <input type="radio"/>	$(S \bowtie T) \bowtie R$
(d) <input checked="" type="radio"/>	$S \bowtie (T \bowtie R)$
(e) <input type="radio"/>	They are all valid join orders for the query.

### Question 3 (1 mark)

Consider the following two relations/tables:

$R(a,b,c)$  where  $r_R = 100000$ ,  $c_R = 100$ ,  $b_R = 1000$

$S(c,d,e)$  where  $r_S = 50000$ ,  $c_S = 500$ ,  $b_S = 100$

If we do a natural join on these two tables, using a block nested loop join with 35 buffers, how many pages do we *read* in completing the join?

(a) <input type="radio"/>	1100
(b) <input type="radio"/>	100000
(c) <input type="radio"/>	None of the other options is correct
(d) <input type="radio"/>	3100
(e) <input checked="" type="radio"/>	4100


### Question 4 (1 mark)

Consider two tables  $R(id,x,y,z)$  and  $S(id,a,b,rid)$  where  $rid$  is a foreign key referring to  $R.id$ . Consider also the following join on these tables:

```
select x,a from R, S where R.id = S.rid;
```

This join is implemented using the hybrid hash join algorithm with 17 memory buffers, 1 buffer used for input (#0), 8 buffers used to hold one partition in memory (#1-#8), 7 output buffers used to transfer tuples in the other partitions to disk (#9-#16), and 1 "spare" output buffer (#16) used only in the first phase (see below).

The diagram below shows the first phase of the hash join, where  $R$  is scanned and partitioned into 8 hash buckets. **Partition 0 is stored in memory; partitions 1-7 are written to disk.**



After the first phase has partitioned  $R$ , the join algorithm then partitions  $S$  using the same hash function. Any tuples that hash to partition 0 are matched against  $R$ 's partition 0 tuples held in memory buffers 1-8, and any resulting matches are written to disk via the spare output buffer. Other  $S$  tuples are written to disk-based partitions. Partitions 1-7 for  $R$  and  $S$  are then processed as for a standard hash join, using a second hash function.

Assume that:

- we have a well-behaved (uniform) hash functions for both the first and second phases
- $R$  contains 3000 tuples in 60 pages; each partition of  $R$  requires exactly 8 pages
- $S$  contains 1600 tuples in 40 pages; each partition of  $S$  requires exactly 5 pages
- each tuple in  $S$  refers to a different tuple in  $R$  (i.e.  $S.rid$  is unique)
- the result contains 1600 tuples which are written into 30 pages (i.e. they count as disk I/O)

Based on the above, compute the number of disk I/Os needed to execute this join.

(a) <input type="radio"/>	308	
(b) <input type="radio"/>	272	$k$ partitions on disk $60 + 40 + 2 * K * (\text{ceil}(\text{br}/ \text{total partitions}) + \text{ceil}(\text{bs}/ \text{total partitions}))$ $= 60 + 40 + 2 * 7 * (8 + 5) = 282$
(c) <input checked="" type="radio"/>	312	plus final write cost + 30 = 312 $60 + 40 + 7 * 8 \text{ writes} + 7 * 5 \text{ reads} + 7 * 8 \text{ reads} + 7 * 5 \text{ writes} + 30 \text{ writes to disk}$
(d) <input type="radio"/>	338	
(e) <input type="radio"/>	None of the other options is correct	

✓ Submit